

# Practical Concurrent and Parallel Programming

Kasper Østerbye

IT University of Copenhagen

Friday 2020-09-04

# Thanks

For using the forum

# Thanks

For using the forum

- for raising questions

# Thanks

For using the forum

- for raising questions
- and for answering them.

# Thanks

For using the forum

- for raising questions
- and for answering them.

The three levels of help:

1. Your group members
2. Your team/timeslot members
3. The forum

# Evaluation of feedback based learning

- I plan to do a questionnaire at the end of September - we need some experience to build on.

# Evaluation of feedback based learning

- I plan to do a questionnaire at the end of September - we need some experience to build on.\* But please mail any concerns to me, or talk with your TA

# Plan for today

- A few notes on yesterdays feedback
- CPU caches and parallelism
- Visibility (of memory writes)
- Instance creation and synchronization



# Exam and the weekly exercises

Written take-home exam (individual mini-project)

- Total duration 30 hours
- Expected exam workload 16 hours
- Individual
- Date and time for exam publication and solution hand-in will be published in LearnIT.
- Submission in LearnIT

# Exam and the weekly exercises

Written take-home exam (individual mini-project)

- Total duration 30 hours
- Expected exam workload 16 hours
- Individual
- Date and time for exam publication and solution hand-in will be published in LearnIT.
- Submission in LearnIT

TIME AND DATE

- Hand out Thu, 10 Dec 2020, 08:00
- Submission Fri, 11 Dec 2020, 14:00

# Exam and the weekly exercises

Written take-home exam (individual mini-project)

- Total duration 30 hours
- Expected exam workload 16 hours
- Individual
- Date and time for exam publication and solution hand-in will be published in LearnIT.
- Submission in LearnIT

TIME AND DATE

- Hand out Thu, 10 Dec 2020, 08:00
- Submission Fri, 11 Dec 2020, 14:00

It is our intention that the exam is "somewhat similar" to the weekly exercises

- Several smaller exercises
- A number of which will be variations of those solved during the semester

# Reentrant

```
class Reentrant <T> {  
    private List<T> myElements;  
  
    public synchronized void add(T element){  
        myElements.add( element ); }  
    public synchronized void addAll(Collection<> col){  
        for (T elem: col) {  
            add( elem ); }  
        }  
    }  
}
```

Which of the following descriptions of synchronized is the correct?

- When a thread try to enter a synchronized method, it will try to obtain the lock. If the lock is already taken, it will wait.
- When a thread try to enter a synchronized method, it will try to obtain the lock. If the lock is already taken by an other thread, it will wait

# Reentrant

```
class Reentrant <T> {  
    private List<T> myElements;  
  
    public synchronized void add(T element){  
        myElements.add( element ); }  
    public synchronized void addAll(Collection<> col){  
        for (T elem: col) {  
            add( elem ); }  
        }  
    }  
}
```

Which of the following descriptions of synchronized is the correct?

- When a thread try to enter a synchronized method, it will try to obtain the lock. If the lock is already taken, it will wait.
- When a thread try to enter a synchronized method, it will try to obtain the lock. If the lock is already taken by an other thread, it will wait

*Reentrancy* is implemented by associating with each lock an acquisition count and an owning thread.

# Why synchronize just to read data?

```
class LongCounter {  
    private long count = 0;  
    public synchronized void increment() {  
        count = count + 1;  
    }  
    public synchronized long get() {  
        return count;  
    }  
}  
// TestLongCounter.java
```

Why needed? The synchronized keyword has **two** effects:

- *Mutual exclusion*: only one thread can hold a lock (execute a synchronized method or block) at a time
- *Visibility of memory writes*: All writes by thread A before releasing a lock (exit synchr) are visible to thread B after acquiring the lock (enter synchr)

# Visibility is really important

```
class MutableInteger { WARNING: Useless
    private int value = 0;
    public void set(int value) { this.value = value; }
    public int get() { return value; }
}
```

- Looks OK, no need for synchronization?
- But thread t may loop forever in this scenario:

```
final MutableInteger mi = new MutableInteger();
Thread t = new Thread(() -> {
    while (mi.get() == 0) { } //Loop while zero
});
t.start();
mi.set(42);
```

This write by thread "main" may be forever invisible to thread t


# Why can the other thread not see the write of 42?

It is because of the caching done inside a modern CPU.

Lets just spend [5 min on this video.](#)



# 42 is just written to main's cache



# Reordering of instructions

```
public class PossibleReordering {  
    static int x = 0, y = 0;  
    static int a = 0, b = 0;  
    public static void main(String[] args) throws InterruptedException {  
        Thread one = new Thread(() -> {  
            a = 1; x = b; });  
        Thread other = new Thread(() {  
            b = 1; y = a; });  
        one.start(); other.start();  
        one.join(); other.join();  
        System.out.println("(" + x + "," + y + ")");  
    }  
}
```

We can get (0,1), (1,0), and (1,1)

# Reordering of instructions

```
public class PossibleReordering {  
    static int x = 0, y = 0;  
    static int a = 0, b = 0;  
    public static void main(String[] args) throws InterruptedException {  
        Thread one = new Thread(() -> {  
            a = 1; x = b; });  
        Thread other = new Thread(() {  
            b = 1; y = a; });  
        one.start(); other.start();  
        one.join(); other.join();  
        System.out.println("( " + x + ", " + y + ")");  
    }  
}
```

We can get (0,1), (1,0), and (1,1)

But also (0,0)

# Reordering of instructions

```
public class PossibleReordering {  
    static int x = 0, y = 0;  
    static int a = 0, b = 0;  
    public static void main(String[] args) throws InterruptedException {  
        Thread one = new Thread(() -> {  
            a = 1; x = b; });  
        Thread other = new Thread(() {  
            b = 1; y = a; });  
        one.start(); other.start();  
        one.join(); other.join();  
        System.out.println("(" + x + "," + y + ")");  
    }  
}
```

We can get (0,1), (1,0), and (1,1)

But also (0,0)

- The Java compiler can *reorder* instructions with no data-dependencies
- Memory caching can make assignments to field-variables appear out of order.

# Aside - why would java reorder statements???

```
a = ...  
b = ...  
//stuff that change a  
//stuff that change b  
a++;  
b++;
```

Sometimes it can yield better performance to let the JVM JIT compiler allow to change this to:

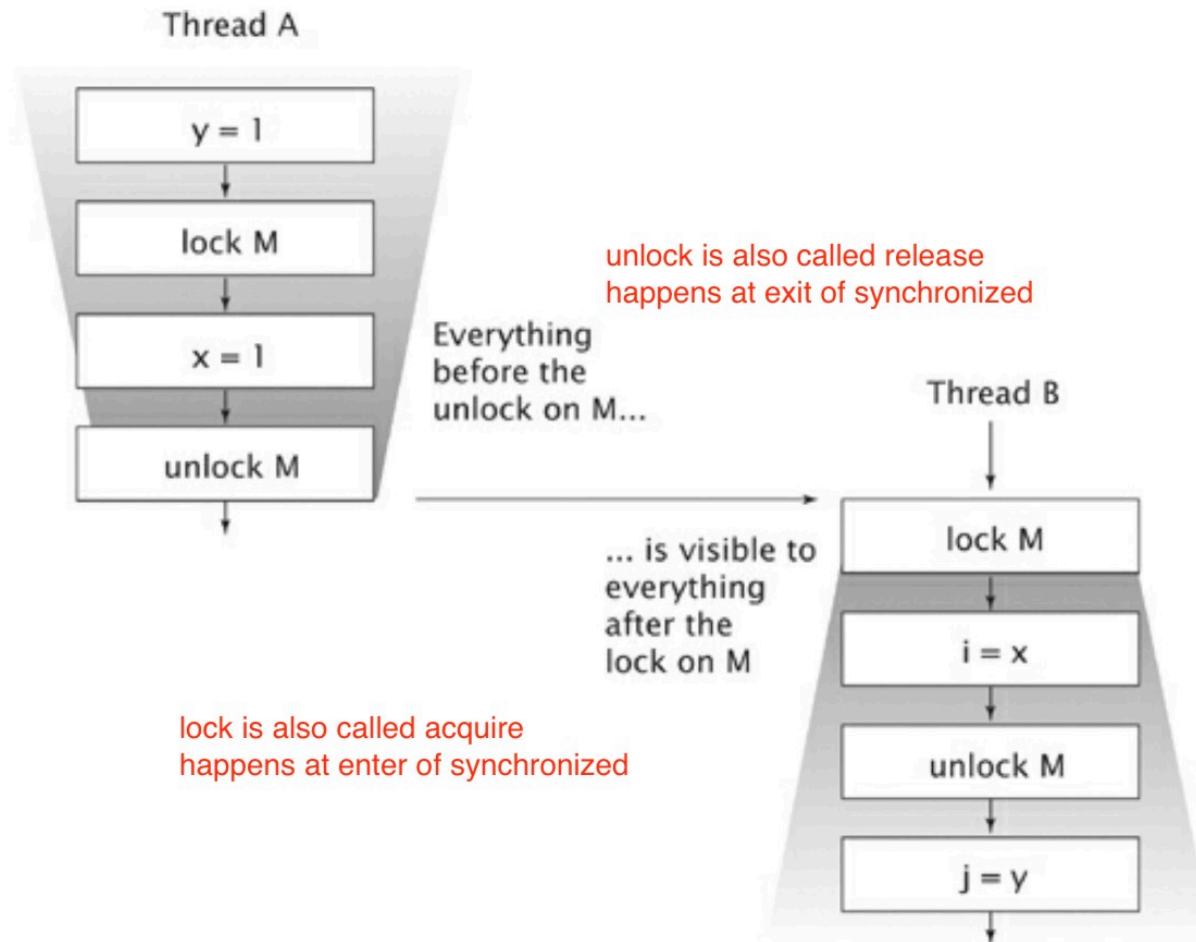
```
a = ...  
//stuff that change a  
a++;  
b = ...  
//stuff that change b  
b++;
```

Because this can allow variable a to remain in one of the cpu registers, and not be written to memory until the end



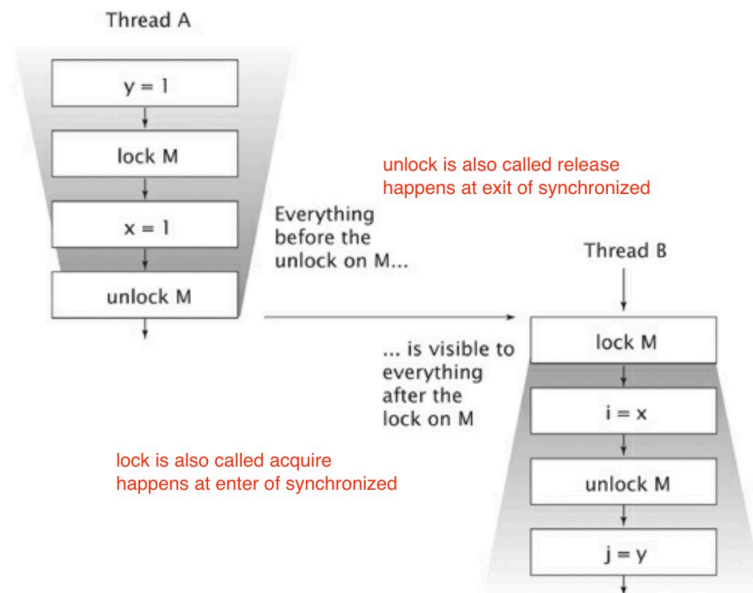
# Visibility by synchronization

One solution is to use synchronized



# MutableInteger - synchronization

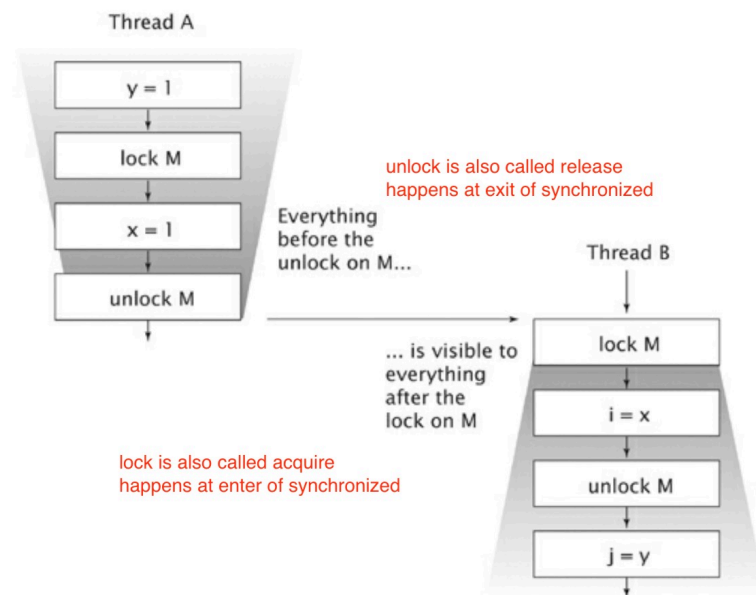
```
class MutableInteger {  
    private int value = 0;  
    public synchronized void set(int value) { this.value = value; }  
    public synchronized int get() { return value; }  
}
```





# MutableInteger - synchronization

```
class MutableInteger {  
    private int value = 0;  
    public synchronized void set(int value) { this.value = value; }  
    public synchronized int get() { return value; }  
}
```



Why is the synchronized keyword needed for get()?

# The volatile field modifier

The volatile field modifier can be used to ensure visibility (but not mutual exclusion)

```
class MutableInteger {  
    private volatile int value = 0;  
    public void set(int value) { this.value = value; }  
    public int get() { return value; }  
}
```

- Volatile variable rule: A write to a volatile field *happens-before* every subsequent read of that same field.

# Goetz advice on volatile

Use volatile variables only when they simplify your synchronization policy; avoid it when verifying correctness would require subtle reasoning about visibility.

Locking can guarantee both visibility and atomicity; volatile variables can only guarantee visibility.

- Rule 1: Use locks (synchronized)
- Rule 2: If circumstances are right, and you are an expert, maybe use volatile instead
- Rule 3: There are few experts

# Ways to ensure visibility

- Unlocking followed by locking the same lock
- Writing a volatile field and then reading it
- Calling one method on a concurrent collection and another method on same collection – `java.util.concurrent.*`
- Calling one method on an atomic variable and then another method on same variable – `java.util.concurrent.atomic.*`
- Finishing a constructor that initializes final or volatile fields Calling `t.start()` before anything in thread `t` Anything in thread `t` before `t.join()` returns

*(Java Language Specification 8 §17.4, and the Javadoc for concurrent collection classes etc, give the full and rather complicated details)*

# Break

# Publishing an object

- What is the problem of making an object, and telling other threads about the object
- Some examples that do not work
- Text book has examples of how it can go wrong

# Unsafe publication

**Listing 16.3. Unsafe Lazy Initialization. *Don't do this.***

```
@NotThreadSafe
public class UnsafeLazyInitialization {
    private static Resource resource;

    public static Resource getInstance() {
        if (resource == null)
            resource = new Resource(); // unsafe publication
        return resource;
    }
}
```



Problems:

- two threads might get into a *race-condition* from the test to the assignment
- Due to caching, you can a Resource which is not initialized (because the fields were not yet written)

# Unsafe publication

**Listing 16.3. Unsafe Lazy Initialization. *Don't do this.***

```
@NotThreadSafe
public class UnsafeLazyInitialization {
    private static Resource resource;

    public static Resource getInstance() {
        if (resource == null)
            resource = new Resource(); // unsafe publication
        return resource;
    }
}
```



Problems:

- two threads might get into a *race-condition* from the test to the assignment
- Due to caching, you can a Resource which is not initialized (because the fields were not yet written)

Would it help to make resource volatile?



# Safe publication

Initialization must *happen-before* publication

If you just want to make it work, use 'synchronized' for `getInstance()`.

# Safe publication

Initialization must *happen-before* publication

If you just want to make it work, use 'synchronized' for `getInstance()`.

Other ways to publish thread-safely:

- Initializing an object reference from a static initializer;
- Storing a reference to it into a volatile field or `AtomicReference`;
- Storing a reference to it into a final field of a properly constructed object;  
or
- Storing a reference to it into a field that is properly guarded by a lock.
- Inserting a reference into a ThreadSafe datastructure.

# Immutable objects

- What are they?
- And how can it go wrong in publishing them too

# Immutable objects

- What are they?
- And how can it go wrong in publishing them too

An object is immutable if:

- Its state cannot be modified after construction;
- All its fields are final and
- It is properly constructed.

# Safe ways to publish objects

We need to assure that initialization *happens-before* publication.

# Safe ways to publish objects

We need to assure that initialization *happens-before* publication.

That is, somewhere the code should be synchronized.

