

# Practical Concurrent and Parallel Programming IV

Jørgen Staunstrup

IT University of Copenhagen

Friday 2020-09-18

Starts at 8:00

# Plan for today

- Performance measurements: motivation and introduction
  - Calculating means and variance (pitfalls)
  - Measurements of thread and lock overhead
    - Questions from you (cost of volatile? and cost of changing thread?)
  - Algorithms for parallel computing
- 

## About Jørgen

Datalog/Computer scientist Aarhus University: 1975

Taught first course on Concurrency: 1978 (USC, Los Angeles)

It was an on-line course !!

Joined ITU in 2001, retired in 2014, Honorary professor since 2017

Teaching MSc course Mobile App Development at ITU since 2016

# Hand-out material

- Peter Sestoft: Microbenchmarks in Java and C sharp
- Slides from lecture
- Exercises week 4
- Recording of lecture (uploaded after the class on 09-18)

## Supplementary material

[Slides from lecture on Computer Arithmetic:](https://www.itu.dk/people/sestoft/bachelor/computernumbers.pdf)

<https://www.itu.dk/people/sestoft/bachelor/computernumbers.pdf>

Take a look at the last page with references (it is a goldmine!!)

## Motivation (performance measurements)

### **Threads are expensive**

How expensive ?

~600 ns to create (on this laptop)

~20 times more time than creating a simple object

40000 ns to start a thread !!!

## Today: How to get such numbers !

# (Performance) Measurements

- Key in many sciences (experiments, observations, predictions, ...)
- A bit of statistics
- A bit of numerical analysis
- A bit of computer architecture (number representation)
- Code for measuring execution time

Based on Microbenchmarks in Java and C# by Peter Sestoft (see [benchmarkingNotes.pdf](#) in material for this week)

All numbers in these slides were measured in September 2020 on a:

Intel(R) Core(TM) i5-1035G4 CPU @ 1.10GHz, 1498 Mhz, 4 Core(s), 8 Logical Processor(s)

# Example

How long does this method take to run?

```
private static double multiply(int i) {  
    double x = 1.1;  
    return x * x * x * x * x * x * x * x * x * x * x * x * x * x * x * x *  
           * x * x * x * x * x * x * x * x * x * x * x * x * x * x * x * x;  
}
```

- 20 floating-point multiplications
- 1 GHz processor ~ 1 ns/cycle
- So takes around 20 ns

2-3 integer operations/cycle (sometimes)

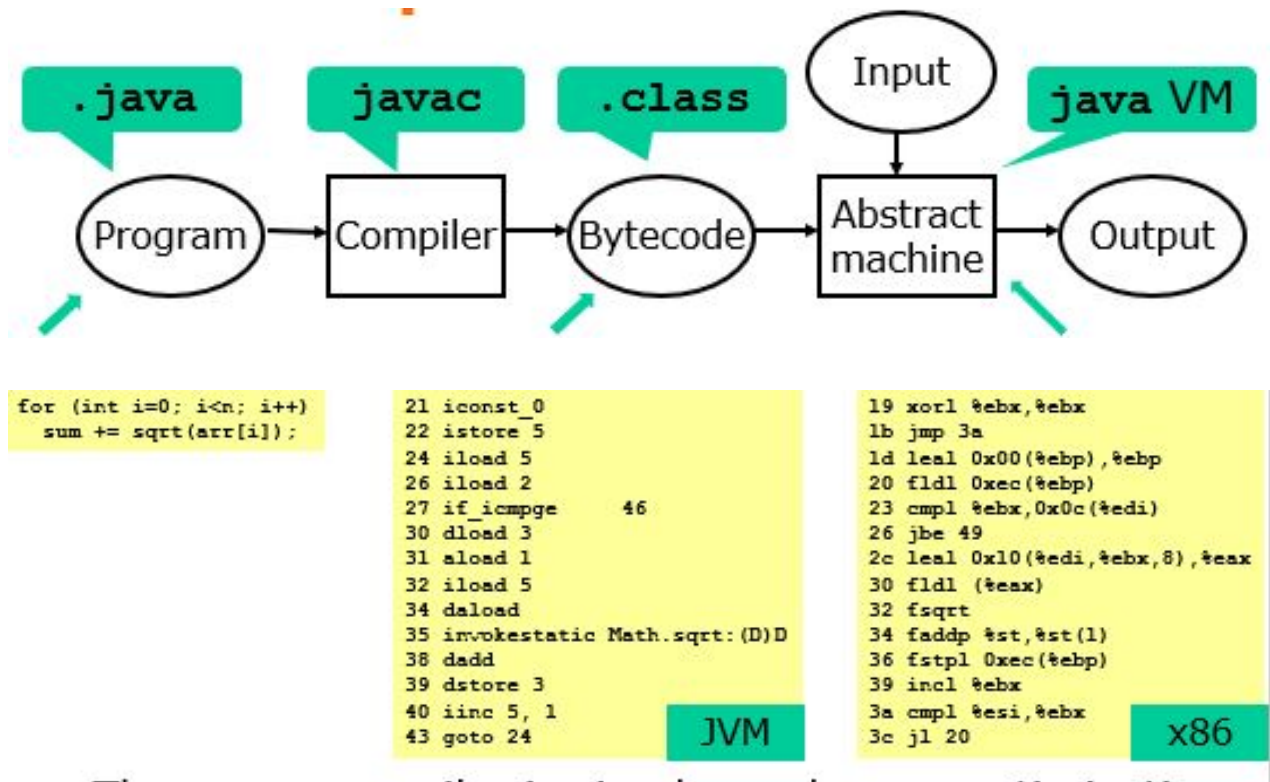
- Let us try to measure it:

```
start = System.nanoTime();  
//calculation  
spent= System.nanoTime()-start;
```

- 1.2 ns ???

JIT optimizes code and removes unused calculations !!!!

# Java compilation and execution



- The javac compiler is simple, makes no optimizations
- The java runtime system (JIT) is clever, obeys spec

"[https://www.tutorialspoint.com/java\\_virtual\\_machine/java\\_virtual\\_machine\\_jit\\_compil](https://www.tutorialspoint.com/java_virtual_machine/java_virtual_machine_jit_compil)"

## Alternative multiplication

```
private static double multiply(int i) {  
    double x = 1.1 * (double)(i & 0xFF);  
    return x * x * x * x * x * x * x * x * x * x * x * x * x * x * x * x *  
           * x * x * x * x * x * x * x * x * x * x * x * x * x * x * x * x;  
}
```

## A simple Timer class for Java

Works on all platforms (Linux, MacOS, Windows)

```
public class Timer {  
    private long start, spent = 0;  
    public Timer() { play(); }  
    public double check()  
    { return (System.nanoTime()-start+spent)/1e9; }  
    public void pause() { spent += System.nanoTime()-start; }  
    public void play() { start = System.nanoTime(); }  
}
```



## Mark2 in BenchMark.java

```
private static double multiply(int i) {  
    double x = 1.1 * (double)(i & 0xFF);  
    return x * x * x * x * x * x * x * x * x * x * x * x * x * x * x * x *  
           * x * x * x * x * x * x * x * x * x * x * x * x * x * x;  
}
```

```
public static double Mark2() {  
    Timer t = new Timer();  
    int count = 100_000_000;  
    double dummy = 0.0;  
    for (int i=0; i<count; i++)  
        dummy += multiply(i);  
    double time = t.check() * 1e9 / count;  
    System.out.printf("%6.1f ns%n", time);  
    return dummy;  
}
```

Let us try it.

## Automating multiple samples (Mark3)

```
int n = 10;
int count = 100_000_000;
double dummy = 0.0;
for (int j=0; j<n; j++) {
    Timer t = new Timer();
    for (int i=0; i<count; i++)
        dummy += multiply(i);
    double time = t.check() * 1e9 / count;
    System.out.printf("%6.1f ns%n", time);
}
```

23.9 ns

23.9 ns

23.7 ns

23.9 ns

23.8 ns

23.7 ns

23.7 ns

23.7 ns

23.7 ns

23.7 ns

# What is the running time?

What should you report as the result, when the observations are:

30.7 ns 30.3 ns 30.1 ns 30.7 ns 30.5 ns 30.4 ns 30.9 ns 30.3 ns 30.5 ns 30.8 ns ??

Mean: 30.4 ns

What if they are:

30.7 ns 100.2 ns 30.1 ns 30.7 ns 20.2 ns 30.4 ns 2.0 ns 30.3 ns 30.5 ns 5.4 ns ??

Mean: 31.0 ns


and variance

# Variance

$$\sigma^2 = \frac{1}{n(n-1)} \left( n \sum_{i=1}^n x_i^2 - \left( \sum_{i=1}^n x_i \right)^2 \right)$$

see also [https://en.wikipedia.org/wiki/Algorithms\\_for\\_calculating\\_variance](https://en.wikipedia.org/wiki/Algorithms_for_calculating_variance)

# Normal Distribution



# What is the running time?

What should you report as the result, when the observations are:

30.7 ns 30.3 ns 30.1 ns 30.7 ns 50.2 ns 30.4 ns 30.9 ns 30.3 ns 30.5 ns 30.8 ns ??

Mean: 32.5 ns Standard deviation: 6.2

50.2 is an outlier

because there is a probability of less than 4.6 % that 50.2 is a correct observation

# Warning

$$\sigma^2 = \frac{1}{n(n-1)} \left( n \sum_{i=1}^n x_i^2 - \left( \sum_{i=1}^n x_i \right)^2 \right)$$

```
int n = 10;
...
for (int j=0; j < n ; j++) {
    Timer t = new Timer();
    for (int i=0; i < count; i++)
        ...
    double time = t.check() * 1e9 / count;
    st += time;
    sst += time * time;
}
double mean = st/n, sdev = Math.sqrt((sst - mean*mean*n)/(n-1));
System.out.printf("%6.1f ns +/- %6.3f%n", mean, sdev);
```

Beware:

`sst - mean * mean * n`

# Floating-point numbers

IEEE 754 binary32 and binary64

Which you know as float and double in Java (and C#)

IEEE 754 decimal128

Java's `java.math.BigDecimal`

The number 1.0527 can be written as:  $10527 / 10^4$  (fraction)

representation is 10527 4

But not all numbers can be represented precisely in a finite number of bits !!!

$1/10 = 0.00011001100110011001100...$

[See Wikipedia page on accuracy problems](#)

So we cannot represent 0.10 Kr. or \$0.10 exactly

Nor 0.01 Kr. or \$0.01 exactly

**Do not use binary floating-point (float, double) for accounting!!!!**



# Digit loss

Beware of cancellation when subtracting numbers that are close to each other:

```
  1010101000010110110001110101.111
- 1010101000010110110001110001.100
-----
  0000000000000000000000000100.011
```

Therefore:

```
(sst - mean*mean)
```

can be problematic.

How to do it:

[https://en.wikipedia.org/wiki/Algorithms\\_for\\_calculating\\_variance](https://en.wikipedia.org/wiki/Algorithms_for_calculating_variance)

## Benchmark code (Mark6)

```
public static double Mark6(String msg, IntToDoubleFunction f) {
    int n = 10, count = 1, totalCount = 0;
    double dummy = 0.0, runningTime = 0.0, st = 0.0, sst = 0.0;
    do {
        count *= 2;
        st = sst = 0.0;
        for (int j=0; j < n; j++ ) {
            Timer t = new Timer();
            for (int i=0; i < count; i++)
                dummy += f.applyAsDouble(i);
            runningTime = t.check();
            double time = runningTime * 1e9 / count;
            st += time;
            sst += time * time;
            totalCount += count;
        }
        double mean = st/n, sdev = Math.sqrt((sst - mean*mean*n)/(n-1));
        System.out.printf("%-25s %15.1f ns %10.2f %10d%n", msg, mean, sdev, count);
    } while (runningTime < 0.25 && count < Integer.MAX_VALUE/2);
    return dummy / totalCount;
}
```

see Benchmark.java (in code for this week)

# Cost of Threads

A thread is an object so let us start finding the cost of creating a simple object.

```
/** @author Brian Goetz and Tim Peierls */  
class Point {  
    public final int x, y;  
    public Point(int x, int y) { this.x = x; this.y = y; }  
}
```

Benchmark code:

```
Mark6("Point creation",  
    i -> {  
        Point p = new Point(i, i);  
        return p.hashCode();  
    });
```

hashCode()	2.6 ns	0.03
Point creation	33.5 ns	0.35

So object creation is: ~ 31 ns

# Thread creation

```
Mark6("Thread create",
    i -> {
        Thread t = new Thread(() -> {
            for (int j=0; j<1000; j++) // not executed
                ai.getAndIncrement(); // thread t created, but not started
        });
        return t.hashCode();
    });
```

Takes 600 ns ~ 20 times slower than Point creation

## Thread create + start

```
Mark6("Thread create start",
    i -> {
        Thread t = new Thread(() -> {
            for (int j=0; j<1000; j++) //most iterations not done
                ai.getAndIncrement(); // Why?
        });
        t.start();
        return t.hashCode();
    });
```

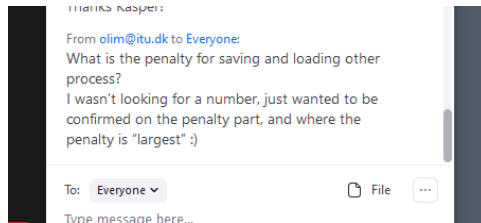
Takes ~ 40000 ns

- So a lot of work goes into setting up a thread
- Even after creating it
- Note: does not include executing the loop (why?)

**Never create threads for small computations**

See more examples in BenchMark.java

# Question from week 3



```
Mark8("Thread yield", i -> { Thread t = new Thread(() -> {  
    for (int j=0; j<100000; j++) Thread.yield();  
});  
t.start();  
try { t.join(); } catch (InterruptedException exn) { }  
return t.hashCode();  
});
```

Also tried the same with and extra Thread t2 to see if Thread.yield(); was optimized away. Apparently, it was not.

Result:

Thread yield	102.2 ns	6.27	2
Thread yield	100.6 ns	1.95	4
Thread yield	100.8 ns	1.26	8
Thread yield	100.6 ns	1.23	16
Thread yield	100.8 ns	0.79	32
Thread yield	100.4 ns	0.75	64
Thread yield	102.9 ns	2.65	128

# Algorithms for parallel computing

## Quicksort

Problem: The first iteration through ALL the data (  $O(n)$  ) only activates one thread, while the rest are waiting.

The second iteration only activates two threads etc.


The resulting speed-up using 8 cores has been reported as 3.6.

## CountPrimes

Problem: Computing the prime factors of small numbers much faster than for large numbers.

Work balancing is IMPORTANT.

# Problem-heap



```
new Thread( () -> {  
    while (newProblem= heap.getProblem()) solveProblem(nextProblem, heap);  
});
```

```
public static void solveProblem(Problem p, ProblemHeap h, LongCounter lc) {  
    int pc= 0; // no of primes  
    if ((p.high - p.low) > threshold) { //generate new problem  
        int d= (p.low+p.high)/2;  
        h.add(new Problem(p.arr, p.low, d));  
        h.add(new Problem(p.arr, d+1, p.high));  
    } else {  
        for (int i=p.low; i<=p.high; i++) if (isPrime(i)) pc= pc+1;  
        lc.incr(pc); // must be atomic  
    }  
}
```



# References

\*Gelernter, David. "Generative communication in Linda". ACM Transactions on Programming Languages and Systems, volume 7, number 1, January 1985

\*Møller-Nielsen, P and Staunstrup, J, Problem-heap. A paradigm for multiprocessor algorithms. Parallel Computing, 4:63-74, 1987

- [Javaspaces](#)

**Problem-heap sorting performance**

**Problem-heap count primes performance**

**Reading and exercises for Week 4**

