

# Practical Concurrent and Parallel Programming VI

Streams & Parallel Streams

IT University of Copenhagen

Friday 2020-10-02

Kasper Østerbye

Starts at 8:00

# Plan for today

Follow up on last week

A bit of history

Streams

Parallel streams

In other languages

# Some history

## Mapping over a list

- Lisp (1958) had lambda expressions and higher order functions. The notion of `map` - applying a lambda expression to each element in a list to get a new list has been around since the beginning of programming.
- Smalltalk (1980) has a large collection library, and also has `map` and `filter` (called `collect` and `select`).

## Lazy eval

Lazy evaluation assigns an *expression* to a variable. When the value of the variable is read, the expression is evaluated.

```
x = 10+y; // corresponds to x = ()-> 10+y.eval()  
println( x ); // corresponds to print( x.eval() ) - printing the value of 10+y
```

The earliest usage is in ALGOL 60. Here it is not used for general variables, but for only parameters. It was named *Jensens Device* after *Jørn Jensen* who worked at Regnecentralen in the early 1960 & 70s.

# More history - lazy streams

The language Icon (1977-2018) was build around lazy streams. All expressions produced stream results.

```
readline = ("bingo" | "banko")
```

This is a boolean expression which is true if any line in input is either "bingo" or "banko".

Once a match is found, no further elements are read, and the streams do not produce any more elements.

## Functional programming

Many functional programming languages use lazy streams. Haskell seems to be one of the earliest to go all in, but most functional programming languages support lazy streams

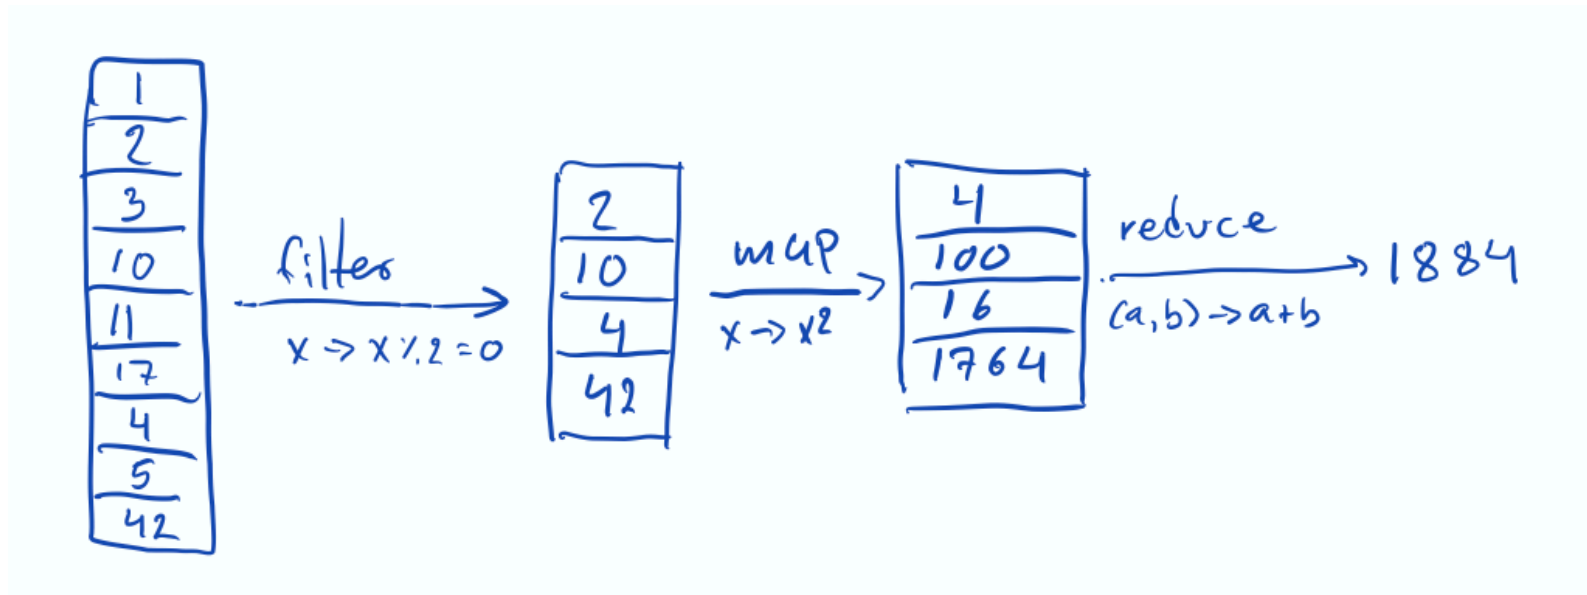
# Mainstream streams

## C# 3.5 (2007)

- First mainstream programming language to support lambda expressions and lazy streams
- It was a revolution, and earned Microsoft a lot of honor in academic circles as it truly is well designed

## Java 8 (2014)


# Filter, map, reduce



## Eager model

- Not supported in Java (but lisp and smalltalk for example)
- New collections are generated at each link in the chain.
- Not used in modern frameworks as it potentially can give many copies of the same data
- Does not support *infinite streams*.


# Lazy or Iterator model



## Lazy/Iterator model

- Used in Java and C# - Python?
- No intermediate collections generated
- Supports *infinite streams*
- Can be difficult to re-start (you cannot (easily) store the chain in a variable and use it several times)

# Stream pipeline



There are three different types of elements:

- sources (arrays, collections, IO, generators)
- intermediate operations (transforming one stream into an other - filter, map,...)
- terminal operations ( count, sum, forEach, ...)

The streams are lazy, driven by the terminal



# Sources

- Input (files or network) - Here from `BufferedReader`

```
Stream<String> lines()  
Returns a Stream, the elements of which are lines read from this  
BufferedReader.
```

- The `Arrays` class has a number of utilities - for example:  
`Arrays.stream(array);`
- All java collections has a `stream()` method which is a stream over its elements
- `Stream.of("Huey", "Dewey" "Louie")` - returns a stream of the three strings  
.

# Example - CSV reader

CSV - comma separated values - used to store tables in text In Denmark we use ; to separate because , is used as decimal character in numbers

Assume we have a file of addresses:

Jomfru Ane gade 17; 9000; Ålborg  
Kannikegade 10; 8000; Aarhus


```
bufferedReader.lines()  
  .map(line -> line.split(";"))  
  .map(elements -> new Address(elements[0], elements[1], elements[2]))  
  .filter( addr -> addr.getZip() >= 9000 )  
  .forEach( Address::toString )
```

# Intermediate operations

- filter - lambda returns a boolean, if the boolean is true the element is included in the output stream
- map - the lambda computes a value. The output stream consists of computed values.
- limit(n) - returns a stream of the first n elements
- skip(n) - returns a stream without the first n elements
- distinct - removes duplicates from stream
- peek - copies the elements to output, but executes lambda for each element (can be used for debugging)
- sorted - returns a stream with the elements sorted

# Terminal operations

- Number streams has min, max, sum, average, count



- Number streams even has `summaryStatistics()`
- `allMatch` - returns true if the lambda is true for all elements - sort of  $\forall$
- `anyMatch` - returns true if the lambda is true for a element () - sort of  $\exists$
- `noneMatch` - returns true if the lambda is false for all elements - sort of  $\nexists$
- `findAny` and `findFirst` returns an `Optional` (as the stream might be empty).
- `forEach` - executes the lambda for all elements in the stream

# Reduce

*Mostly you will not need this, but use one of the predefined ones (sum, average)*

```
stream.reduce( (a,b) -> a+b*b)
```

will compute the sum of squares

the result is computed like this:

```
boolean foundAny = false;
T result = null;
for (T element : this stream) {
    if (!foundAny) {
        foundAny = true;
        result = element;
    }
    else
        result = theLambda.apply(result, element);
}
return foundAny ? Optional.of(result) : Optional.empty();
```

Notice: It returns an Optional as the stream might be empty

# Reduce with initial value

```
stream.reduce(0, (a,b) -> a+b*b)
```

which is computed as:

```
T result = identity;  
for (T element : this stream)  
    result = theLambda.apply(result, element)  
return result;
```

Notice: No optional - as we give an initial value.

# Java type in reduce

That is, one **cannot** do this:

```
List<Integer> list = Arrays.asList(1,2,3,10,11,17,4,5,42);  
System.out.println("Reduced to: " + list.stream()  
    .filter( x -> x%2 == 0)  
    .map( x -> x*x )  
    .reduce(LocalDate.now(), (a,b) -> a.plusDays( b ) ));
```

Which to me looks quite normal....

# Stream to array

Often one want to produce a collection of the elements in a stream.

- `toArray()` - returns an `Object[]` - that is, untyped elements
- `toArray(IntFunction<A[]> generator)` - returns an array of type A The parameter looks a bit odd, but what is ask for is a lambda which takes a size and returns an array of type A of that size.

```
intStream.filter( i -> i.isPrime()).toArray( n -> new Integer[n]);
```

which can be written as

```
intStream.filter( i -> i.isPrime()).toArray(Integer[]::new);
```



# Stream to Lists

- `toList()` - yes, finally something simple

# Stream to ...

There is a general mechanism:

```
stream.collect(aCollector)
```

The java class [Collectors](#) has many different factory methods for aggregating the elements in a stream.

The class has several examples as well.


# Reduce and collect

"The normal reduction is meant to **combine two immutable values** such as int, double, etc. and produce a new one; it's an immutable reduction. In contrast, the **collect method is designed to mutate a container to accumulate the result** it's supposed to produce."

It is possible to write your own collectors, but in nearly all cases you can depend on those declared in

[Collectors](https://docs.oracle.com/javase/8/docs/api/java/util/stream/Collectors.html)(<https://docs.oracle.com/javase/8/docs/api/java/util/stream/Collectors.html>)

# Grouping - principle



each group is represented by a value of the lambda, and all the elements which gave that value

(Sorry for the round stream, was just easier to draw)

The full grouping is a Map.

$\text{Map}<K, \text{List}<T>>$ , where  $K$  is the type of the results of the lambda, and  $T$  is the type of the values in the stream.

# Grouping example

```
Map<Integer,List<Integer>> groupings = Stream.of(1,2,3,10,11,17,4,5,42)
    .collect(Collectors.groupingBy(i-> i / 5)); // integer division by 5


for (int key: groupings.keySet()){
    System.out.format("Key: %d with list: %s\n",
        key, groupings.get(key).toString());
}
```

Prints:

```
Key: 0 with list: [1, 2, 3, 4]
Key: 1 with list: [5]
Key: 2 with list: [10, 11]
Key: 3 with list: [17]
Key: 8 with list: [42]
```

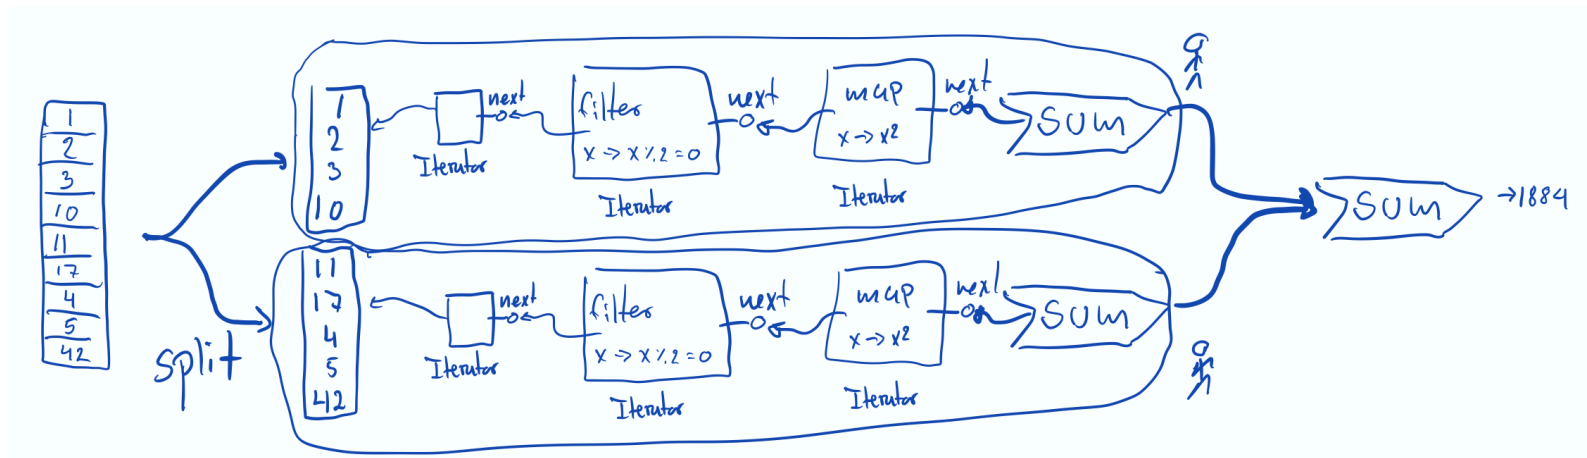
# Parallelism and streams I

**Strategy 1 - parallel pipelining - (not used):**



# Parallelism and streams II

Strategy 2 - splitting the iterator - (used):



# In java

```
anyStream.parallelStream().allAsBeforeStillWorksButIsNowDoneInParallel
```

A special case is with grouping:

```
Map<Person.Sex, List<Person>> byGender =  
    personStream  
        .collect(  
            Collectors.groupingBy(Person::getGender));
```

The following is the parallel equivalent:

```
ConcurrentMap<Person.Sex, List<Person>> byGender =  
    personStream  
        .parallelStream()  
        .collect(  
            Collectors.groupingByConcurrent(Person::getGender));
```



# C

- Java streams are pretty much modeled after C#
- C# has a different kind of type parameters - allowing them to have streams of primitive types
- C# has *extension methods*, which in many cases gives a more concise syntax
- C# has *anonymous classes*, which work well with `map` (called *select* in C#)

# Mongo aggregation

- the source is collection of json documents
- filter allows to look down into the json tree
- map allow to pick out subtrees and substitute elements
- sum, average,... all exists
- a special unwind operator can create copies of a json node  
for example, if a person has three emails  
emails: ["email1", "email2", "email3"]  
unwind will produce three persons, each with only one email.

It is a bit unclear how parallelism is used except on large scale (sharding).

# Python, Ruby, Go, Rust, Kotlin, Scala, F#,...## Left as an exercise

- The devil is in detail

