

# Practical Concurrent and Parallel Programming X

## Coroutines II

**IT University of Copenhagen**

**Friday 2020-11-06**

**Kasper Østerbye**

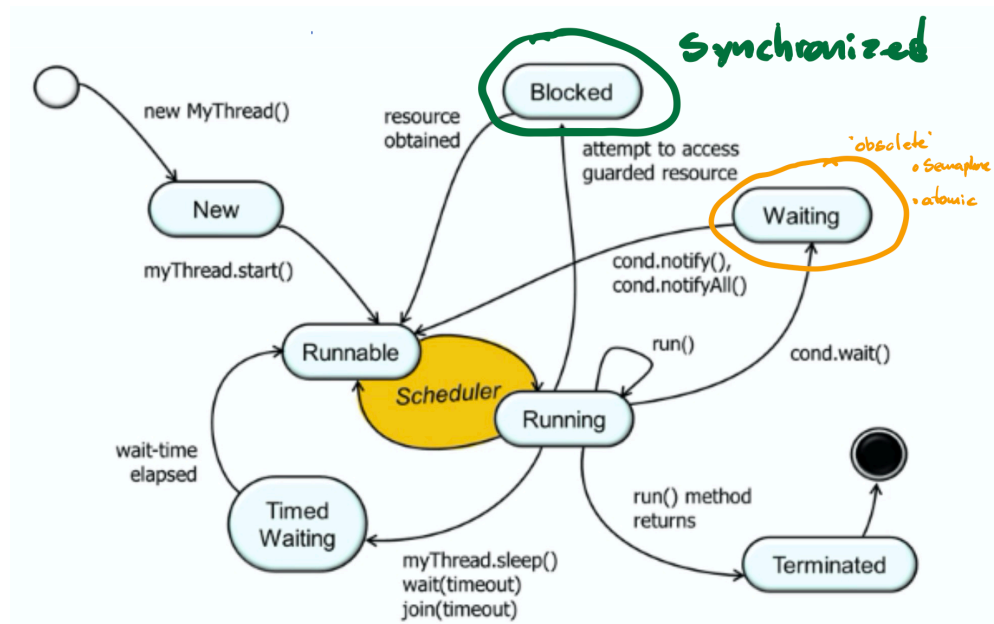
**Starts at 8:00**

# Plan for today

Follow up on last week

# Follow up and loose ends

# Scheduler of Java threads



# Recap

The java virtual machine

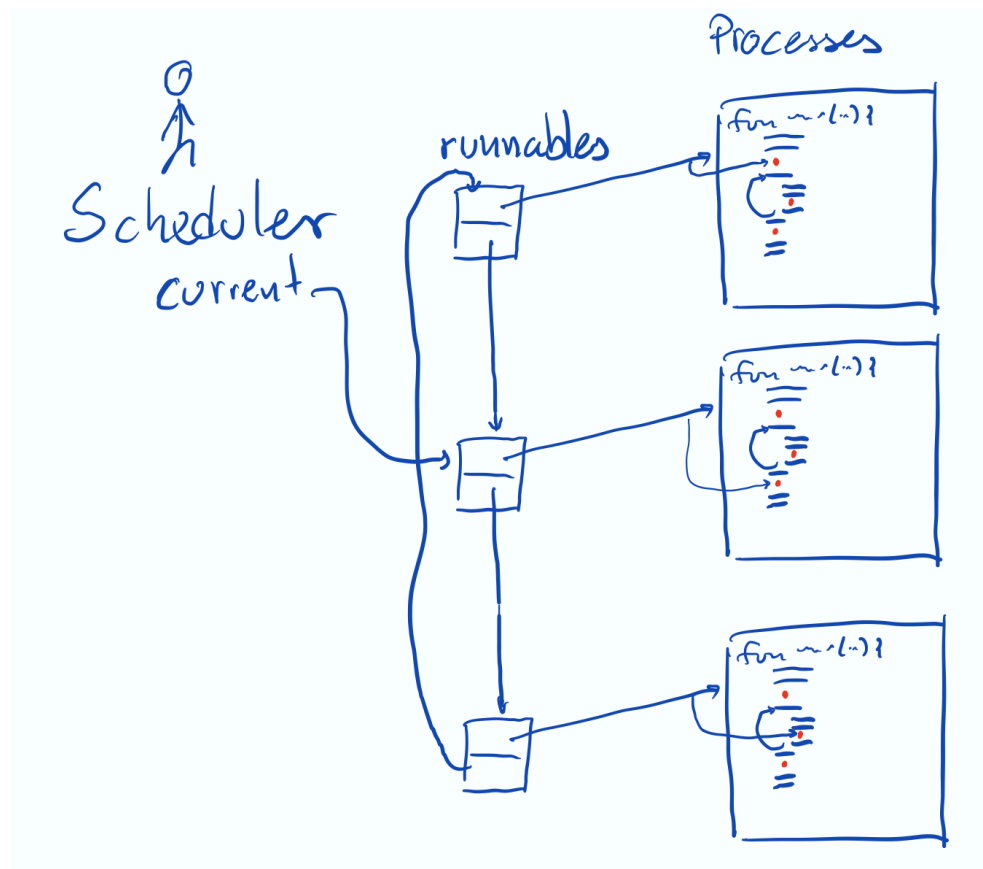
1. starts the sheduler
2. creates a (main) thread
3. which becomes runnable
4. and the scheduler runs it

The JVM scheduler maintains:

1. a list of all threads
2. a low-level mechanism which can put run a thread
3. policies for when to "pre-empt" a running thread and move it to "runnable"

The JVM scheduler runs until the jvm is terminated

# Data structures in a Scheduler



# Scheduler algorithm

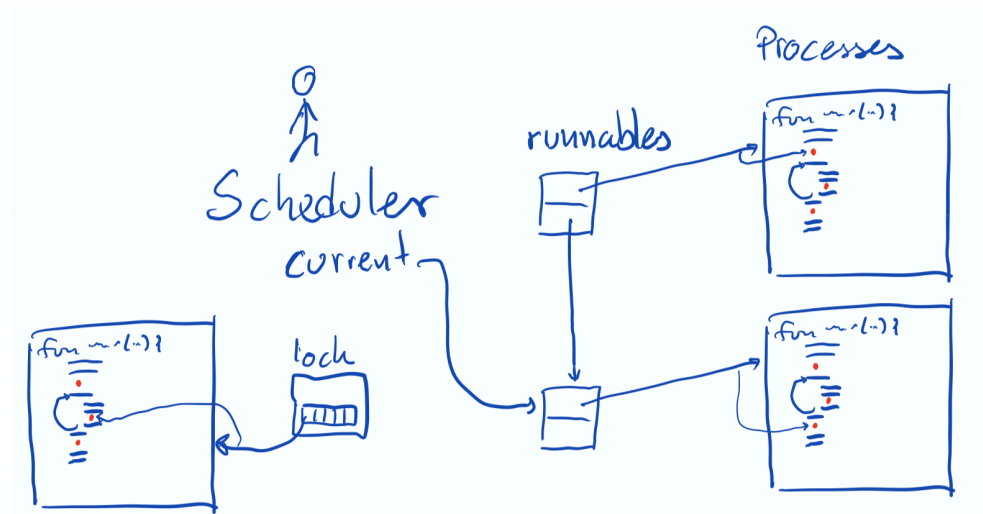
```
void Scheduler.run(){
    current := runnables.first()
    while (true) {
        execute(current, 50ms)
        current := runnables.next()
    }
}
```

This is *pseudo code*.

The execute( current, 50ms ) will

- terminate after 50ms,
- or if the process itself pauses (locks, delays,...)

# With locking



Notice, it is the Processes themselves which move in and out of the locks.

The fundamental objects like the runnable list, the locks, the processes, and scheduler is **designed together**



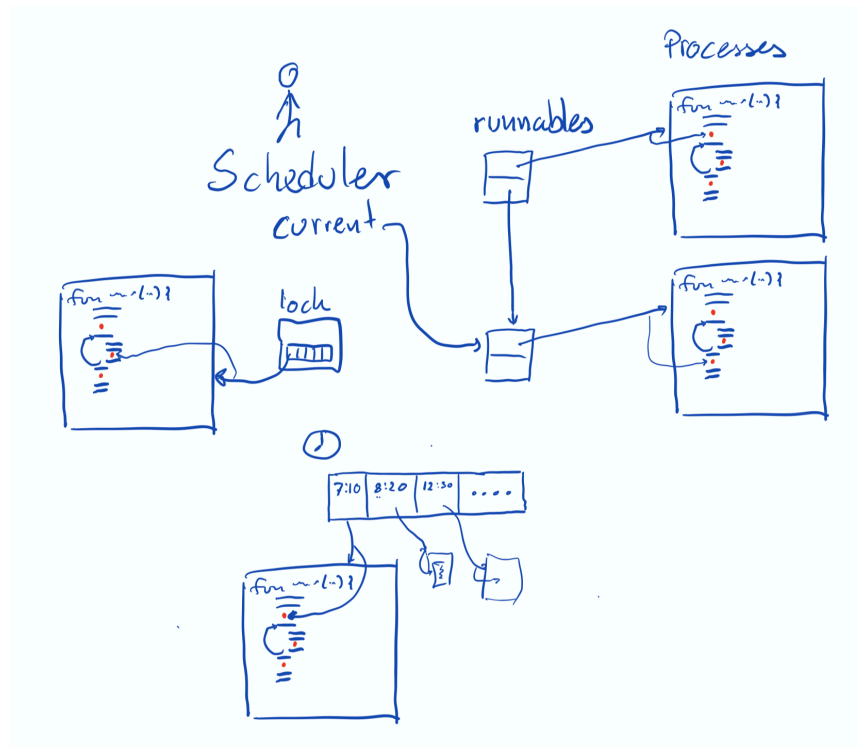
# Locking code

```
// In some Process
myLock.lock()
// do safe stuff
myLock.release()
...
```

```
atomic void Lock.lock(){
    if ( this.closed() ) {
        this.queue.insert( Scheduler.current )
        Scheduler.runnables.remove( Scheduler.current )
        Scheduler.Pause()
    }
    this.close()
}
```

```
atomic void Lock.release(){
    if ( !this.lock.queue.empty() ){
        Scheduler.runnables.add( this.lock.queue.removeFirst() )
    }
}
```

# With locking and timing



It is the current process itself which calls sleep.

# Timing code

```
atomic void Process.sleep(int millies){
    timerQueue.insert(this, currentTime+millies)
    Scheduler.runnables.remove( Scheduler.current )
    Scheduler.pause()
}
```

The scheduler loop now looks like:

```
atomic void Scheduler.run(){
    current := runnables.first()
    while (true) {
        execute(current, 50ms)
        if (currentTime >= timerQueue.first.time){
            current = timerQueue.removeFirst()
            Scheduler.runnables.add( current )
        } else {
            current := runnables.next()
        }
    }
}
```


# Why locking?

To ensure safe access to shared data (scheduling & visibility)

Any alternatives?

Indeed

# Message based concurrency




- Each process has its own state
- Can send/receive immutable messages on channels
- Can wait for messages on one or more channels
- (In general it will be rare that all processes are connected to all others)

The **go programming language** use this as their model for concurrency/parallelism

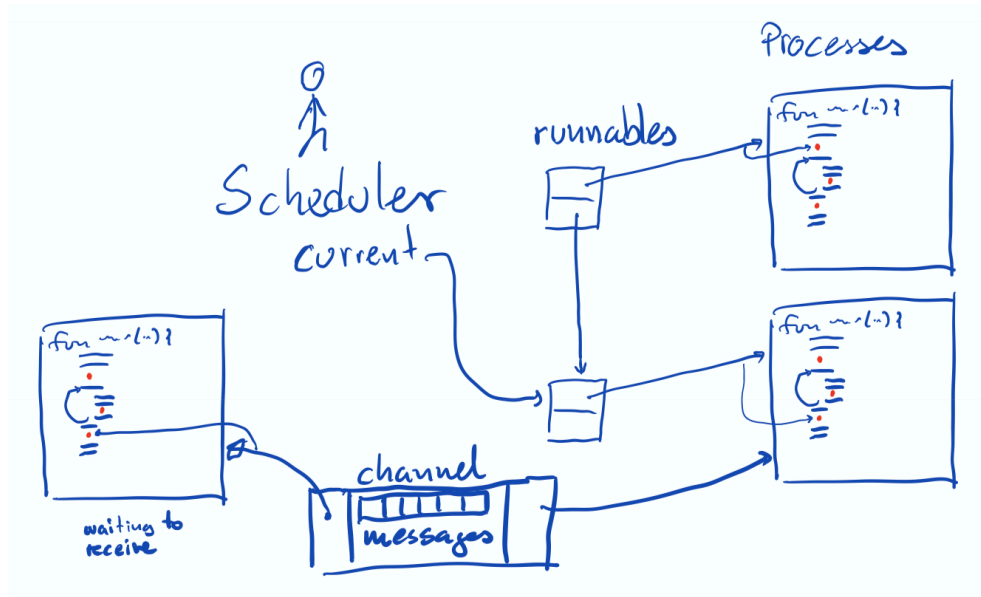
The idea trace back to [Communicating sequential processes by Tony Hoare in 1978](#).

# Message based variations



- a channel *might be* multiway (connect more than two processes)
  - a sent message *might be* received one or by all
- a channel *might have* direction
- a channel *might have* capacity
- waiting *might have* timeout

# Message based scheduler



The left process is not runnable as it is waiting for some other process to send a message on the channel.

(There is only drawn one channel, many could exist. The channel is depicted as being able to store multiple messages, there is no explicit direction on this channel, no clock is associated with this wait)

# Channels in kotlin

```
fun main() = runBlocking {  
    val channel = Channel<Int>()  
    launch { senderRoutine( channel ) }  
    launch { receiverRoutine( channel ) }  
    println("Coroutines started")  
}
```



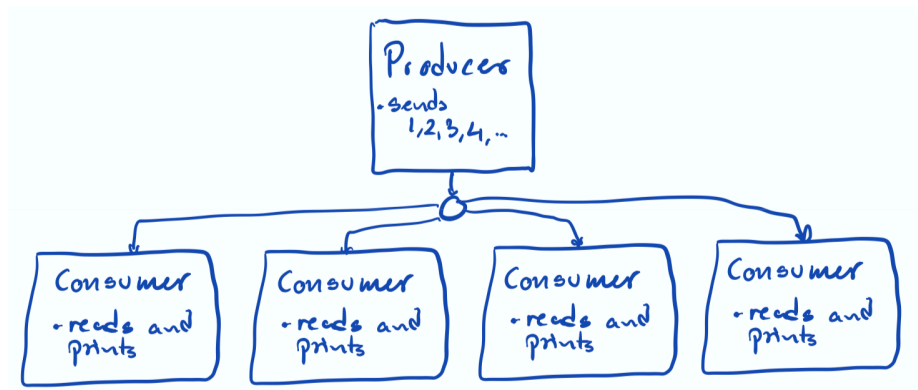


## The [code for the two coroutines](#)

```
suspend fun senderRoutine(channel: Channel<Int>) {  
    for (x in 1..5) {  
        channel.send( x * x )  
        println("Sent ${x*x}")  
    }  
    channel.send(-1) // we're done sending  
    println("Done sending")  
}
```

```
suspend fun receiverRoutine(channel: Channel<Int>) {  
    while(true) {  
        val square = channel.receive()  
        if (square == -1) break  
        println( "Received: $square" )  
    }  
    println("Recieved all!")  
}
```

# Fan-out



In the kotlin implementation, only one consumer will receive a sent message.

```
fun main() = runBlocking<Unit> {  
    val channel = Channel<Int>()  
    launch { produceNumbers( channel ) }  
    repeat(5) { launch { consumer(it, channel) } }  
    delay(1000) // main coroutine is waiting a second  
    channel.cancel()  
}
```

## Code for producer and consumer

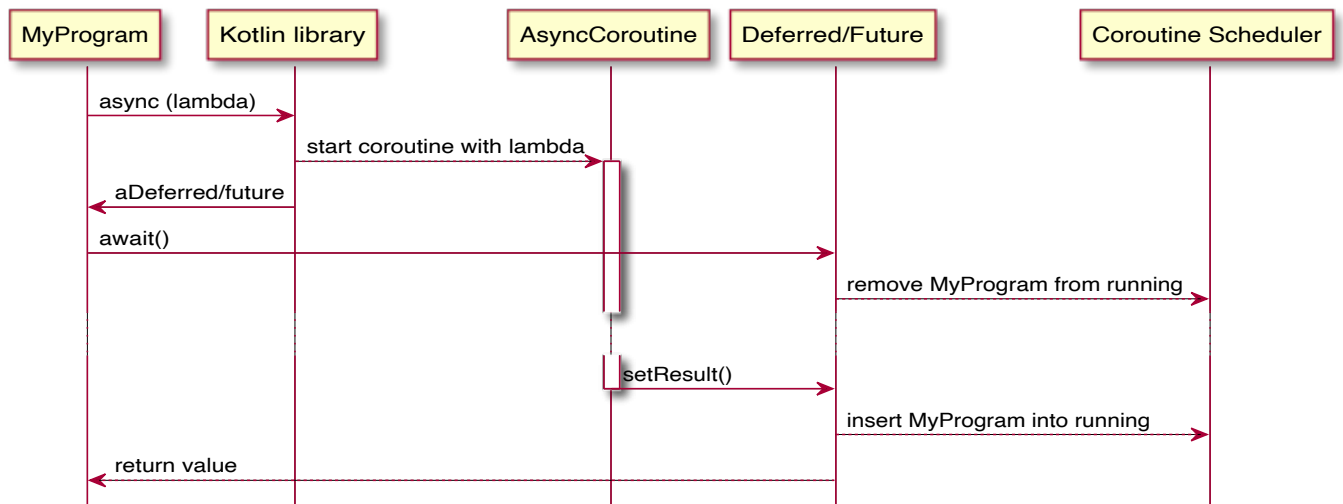
```
suspend fun produceNumbers(ch : Channel<Int>) {  
    var x = 1 // start from 1  
    while (true) {  
        ch.send(x++) // produce next  
        delay(100) // wait 0.1s  
    }  
}
```

```
suspend fun consumer(id: Int, channel: ReceiveChannel<Int>) {  
    for (msg in channel) { // reading from a channel  
        println("#$id --> $msg")  
    }  
}
```

# Async & await

```
val one = async { computation1() }  
val two = async { computation2() }  
res = one.await() + two.await()
```

- The function `async` launches a coroutine,
- returns a `Deferred` (Kotlin name for future/promise)
- The `await()` method on a deferred waits for the result to be ready



# Library defined scheduler

In some languages it is possible to write a scheduler in the language itself.

```
Scheduler sch = new Scheduler()  
sch.add( new Process() )  
sch.add( new Process() )  
sch.run()
```

Where Scheduler and Process are library defined classes (opposed to compiler defined).

# Design options

1. Should `run` return a value
2. How should `Scheduler` work together with native `Threads`
3. Can you add `Process` after the scheduler has started?
4. Which *states* do a process have (running, runnable, waiting, terminated,...)
5. How do you specify a `Process` (using `Lambda` or subclassing for example)
6. Can `Scheduler` be preemptive, or do `Process` have to be cooperative

# Mapping coroutines to JVM threads

```
fun main() = runBlocking<Unit> {  
    launch { println("Parent") }  
    launch(Dispatchers.Unconfined) { println("Unconfined") }  
    launch(Dispatchers.Default) { println("Default") }  
    launch(newSingleThreadContext("MyOwnThread")) { println("New") }  
}
```

Dispatchers are schedulers that run coroutines.

- `launch()` - context of the parent, main `runBlocking` coroutine
- `launch(Dispatchers.Unconfined)` - main thread
- `launch(Dispatchers.Default)` - a `CachedThreadPool`
- `launch(newSingleThreadContext("MyOwnThread"))` - will get its own new thread

[Link to code](#)

# Move coroutines between schedulers

```
fun log(msg: String) = println("[${Thread.currentThread().name}] $msg")

fun main() {
    val ctx1 = newSingleThreadContext("Ctx1")
    val ctx2 = newSingleThreadContext("Ctx2")
    runBlocking(ctx1) {
        log("Started in ctx1")
        withContext(ctx2) {
            log("Working in ctx2")
        }
        log("Back to ctx1")
    }
}
```

[Code](#)

OK, it is possible, but can that have any possible use????



# Thread confinement of access to shared data

**No need to sync if only one thread modifies data**

If coroutines are run in thread-pool with true parallelism, bad things™ might happen.

**The problem - in kotlin, but we saw this in first lesson**

Sample from [documentation](#)

# The fastest solution

(because this has an easy solution)

# The sheduling solution

# A final word on yield

When the scheduler is an object (as opposed to a JVM concept), one can call methods on a scheduler.

The sequence function (where we use yield) is actually a scheduler with just one Process

