

Practical Concurrent and Parallel Programming XI

Coroutines II

IT University of Copenhagen

Friday 2020-11-06

Kasper Østerbye

Starts at 8:00

Plan for today

Follow up on last week

Scheduling

Channels

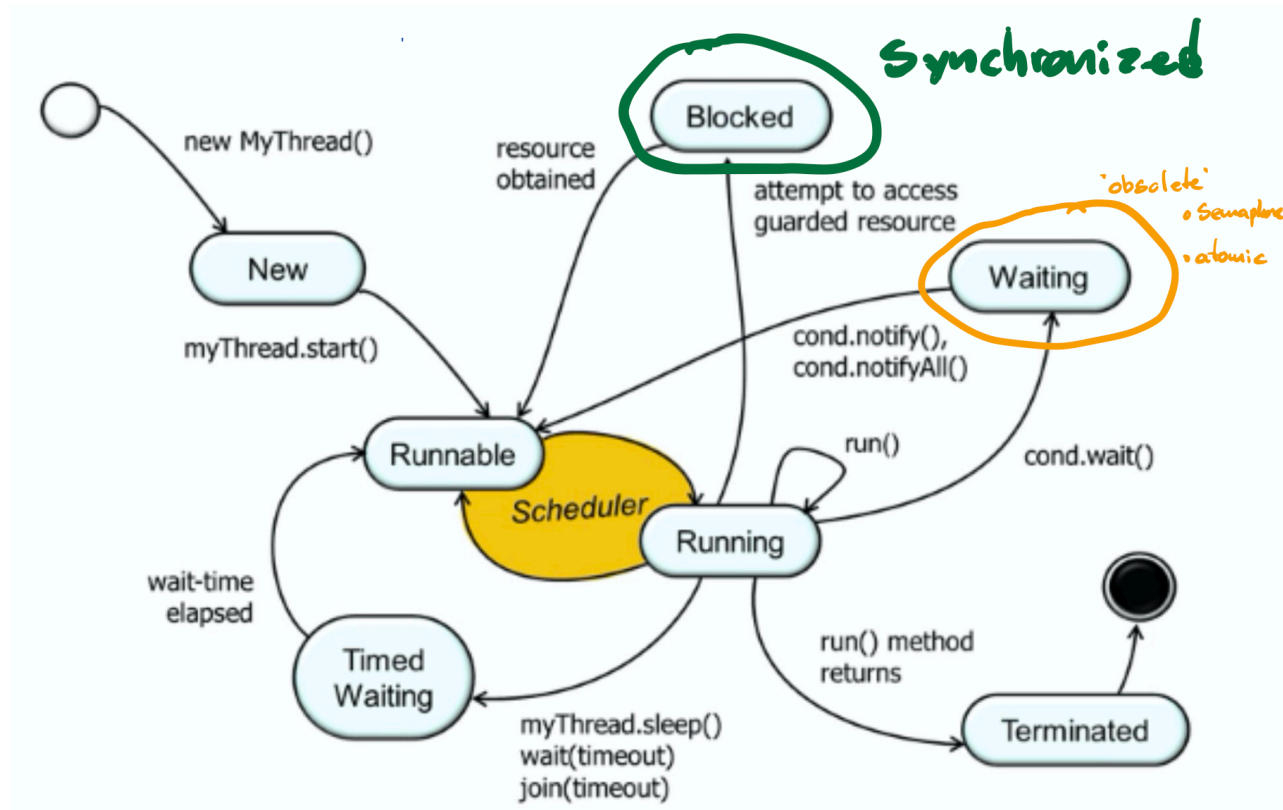
There was a number of questions regarding the statemachine exercise. I would like to postpone this to the end of todays lecture

Section - Schedulers

The scheduler is the piece of code which enables several threads to run on the same core.

Any system with more threads than cores needs a scheduler.

Scheduler of Java threads



Recab

The java virtual machine

1. starts the sheduler
2. creates a (main) thread
3. which becomes runnable
4. and the scheduler runs it

The JVM scheduler maintains:

1. a list of all threads
2. a low-level mechanism which can run a thread
3. policies for when to "pre-empt" a running thread and move it to "runnable"

The JVM scheduler runs until the jvm is terminated

Data structures in a Scheduler

Scheduler algorithm

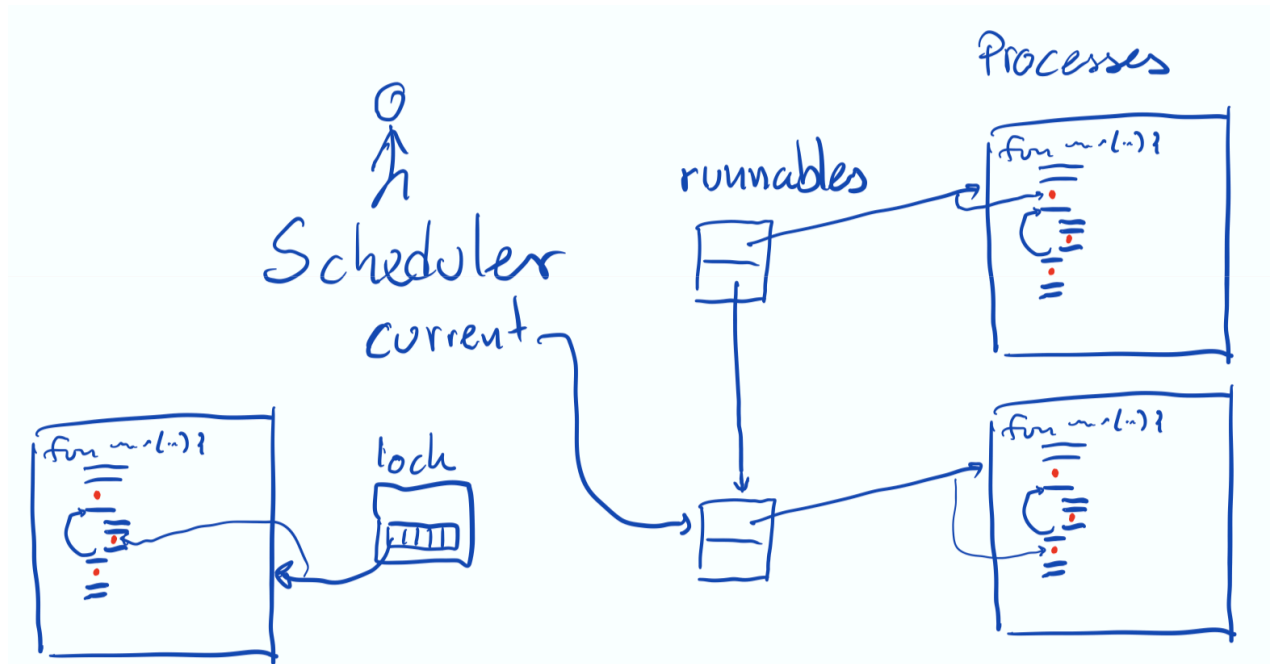
```
void Scheduler.run(){
    current := runnables.first()
    while (true) {
        execute(current, 50ms)
        current := runnables.next()
    }
}
```

This is *pseudo code*.

The execute(current, 50ms) will

- terminate after 50ms,
- or if the process itself pauses (locks, delays,...)

With locking



Notice, it is the Processes themselves which move in and out of the locks.

The fundamental objects like the *runnable* list, the *locks*, the *processes*, and *scheduler* are **designed together**

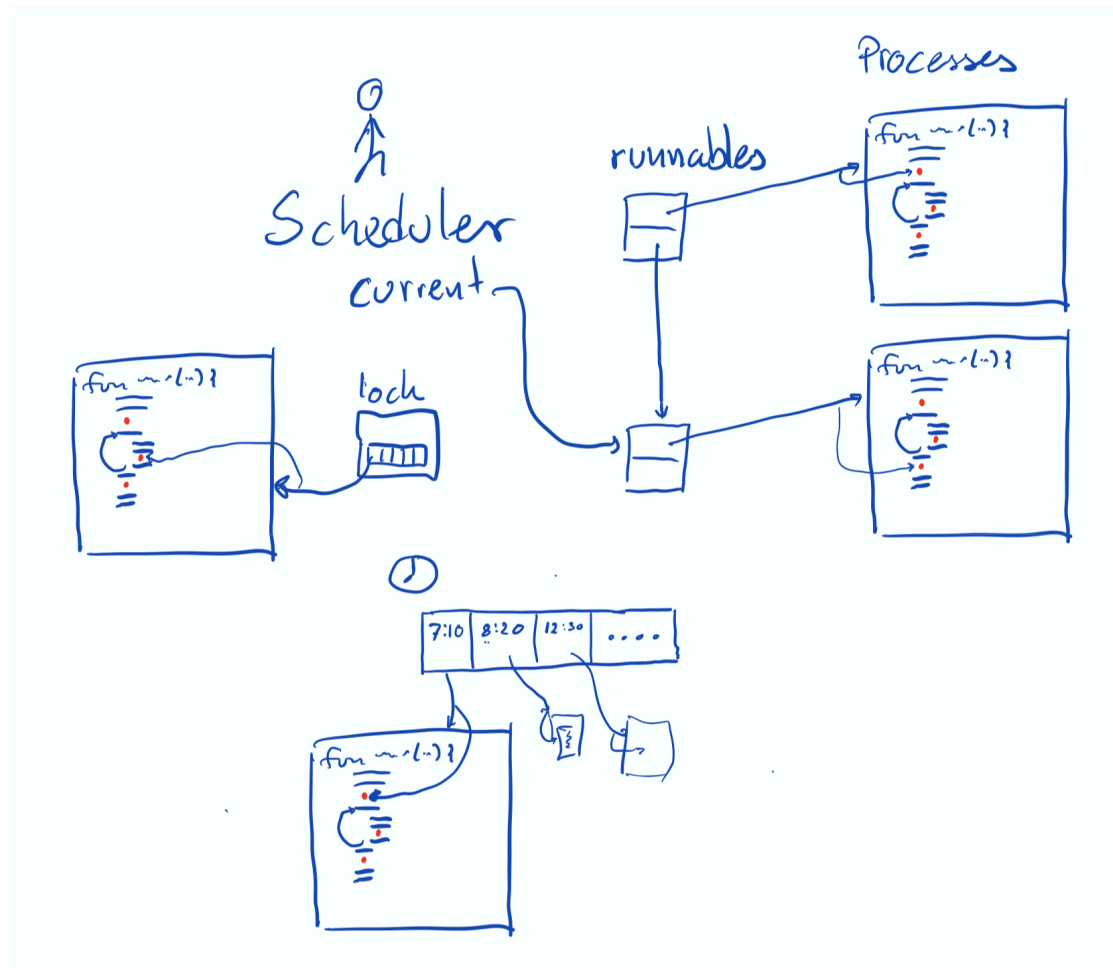
Locking code

```
// In some Process  
myLock.lock()  
// do safe stuff  
myLock.release()  
...
```

```
atomic void Lock.lock(){  
    if ( this.closed() ) {  
        this.queue.insert( Scheduler.current )  
        Scheduler.runnables.remove( Scheduler.current )  
        Scheduler.Pause()  
    } else {  
        this.close()  
    }  
}
```

```
atomic void Lock.release(){  
    if ( !this.lock.queue.empty() ){  
        Scheduler.runnables.add( this.lock.queue.removeFirst() )  
    } else {  
        this.open()  
    }  
}
```

With locking and timing



It is the current process itself which calls sleep.

Timing code

```
atomic void Process.sleep(int millies){
    timerQueue.insert(this, currentTime + millies)
    Scheduler.runnables.remove( Scheduler.current )
    Scheduler.pause()
}
```

The scheduler loop now looks like:

```
atomic void Scheduler.run(){
    current := runnables.first()
    while (true) {
        execute(current, 50ms)
        if (currentTime >= timerQueue.first.time){
            current = timerQueue.removeFirst()
            Scheduler.runnables.add( current )
        }
        current := runnables.next()
    }
}
```


Why locking?

To ensure safe access to shared data (scheduling & visibility)

Any alternatives?

Indeed

Message based concurrency




- Each process has its own state
- Can send/receive immutable messages on channels
- Can wait for messages on one or more channels
- (In general it will be rare that all processes are connected to all others)

The **go programming language** use this as their model for concurrency/parallelism

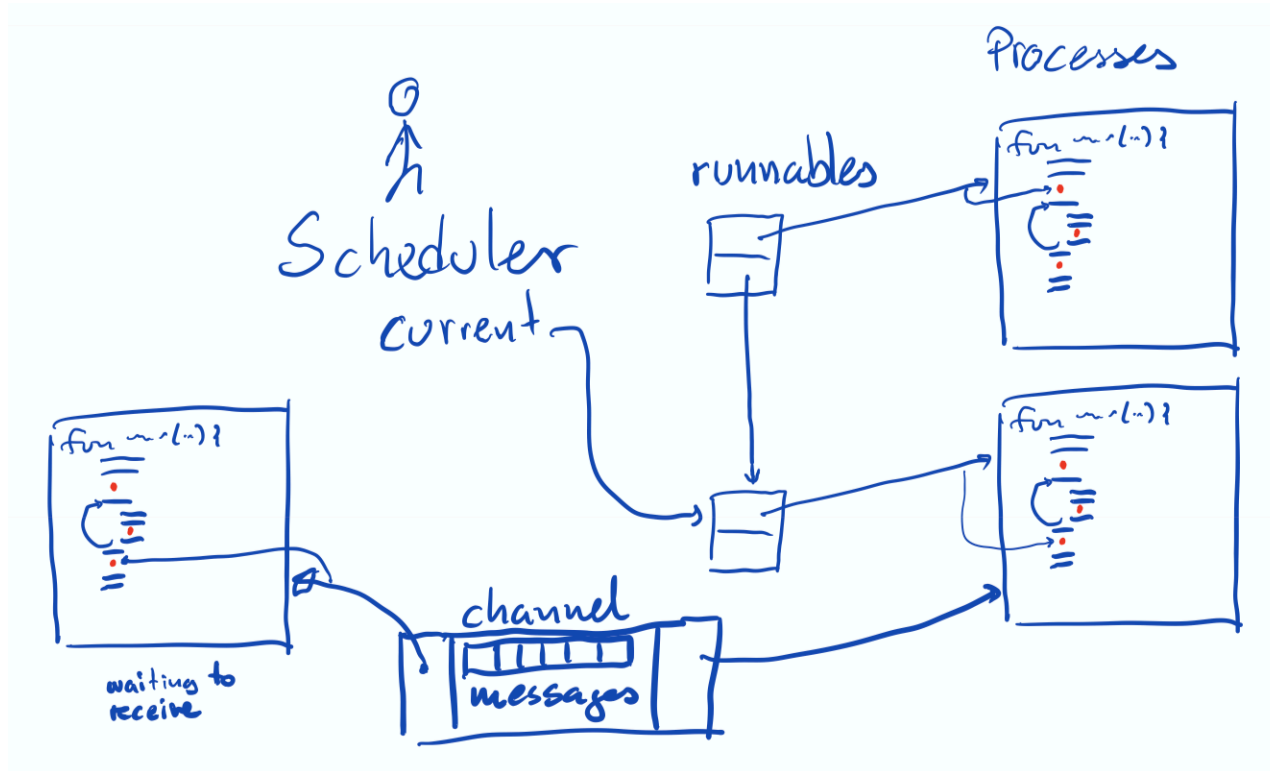
The idea trace back to [Communicating sequential processes by Tony Hoare in 1978.](#)

Message based variations



- a channel *might be* multiway (connect more than two processes)
 - a sent message *might be* received one or by all
- a channel *might have* direction
- a channel *might have* capacity
- waiting *might have* timeout

Message based scheduler




The left process is not runnable as it is waiting for some other process to send a message on the channel.

(There is only drawn one channel, many could exist. The channel is depicted as being able to store multiple messages, there is no explicit direction on this

Channels in kotlin

```
fun main() = runBlocking {  
    val channel = Channel<Int>()  
    launch { senderRoutine( channel ) }  
    launch { receiverRoutine( channel ) }  
    println("Coroutines started")  
}
```



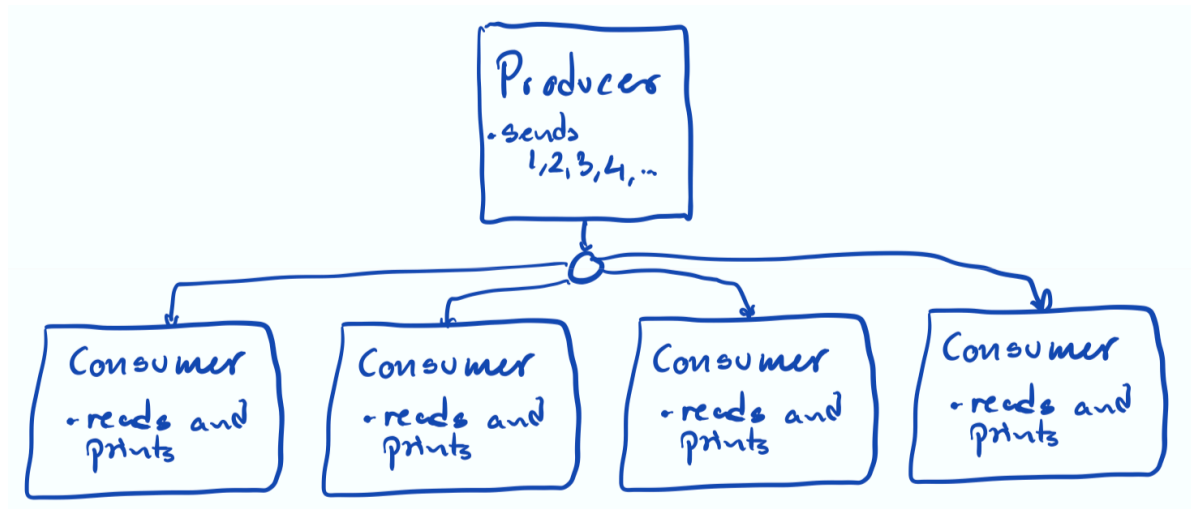
The code for the two coroutines

```
suspend fun senderRoutine(channel: Channel<Int>) {  
    for (x in 1..5) {  
        channel.send( x * x )  
        println("Sent ${x*x}")  
    }  
    channel.send(-1) // we're done sending  
    println("Done sending")  
}
```

```
suspend fun receiverRoutine(channel: Channel<Int>) {  
    while(true) {  
        val square = channel.receive()  
        if (square == -1) break  
        println( "Received: $square" )  
    }  
    println("Recieved all!")  
}
```

Break until 9:00

Fan-out



In the kotlin implementation, only one consumer will receive a sent message.

```
fun main() = runBlocking<Unit> {  
    val channel = Channel<Int>()  
    launch { produceNumbers( channel ) }  
    repeat(5) { launch { consumer(it, channel) } }  
    delay(1000) // main coroutine is waiting a second  
    channel.cancel()  
}
```

Code for producer and consumer

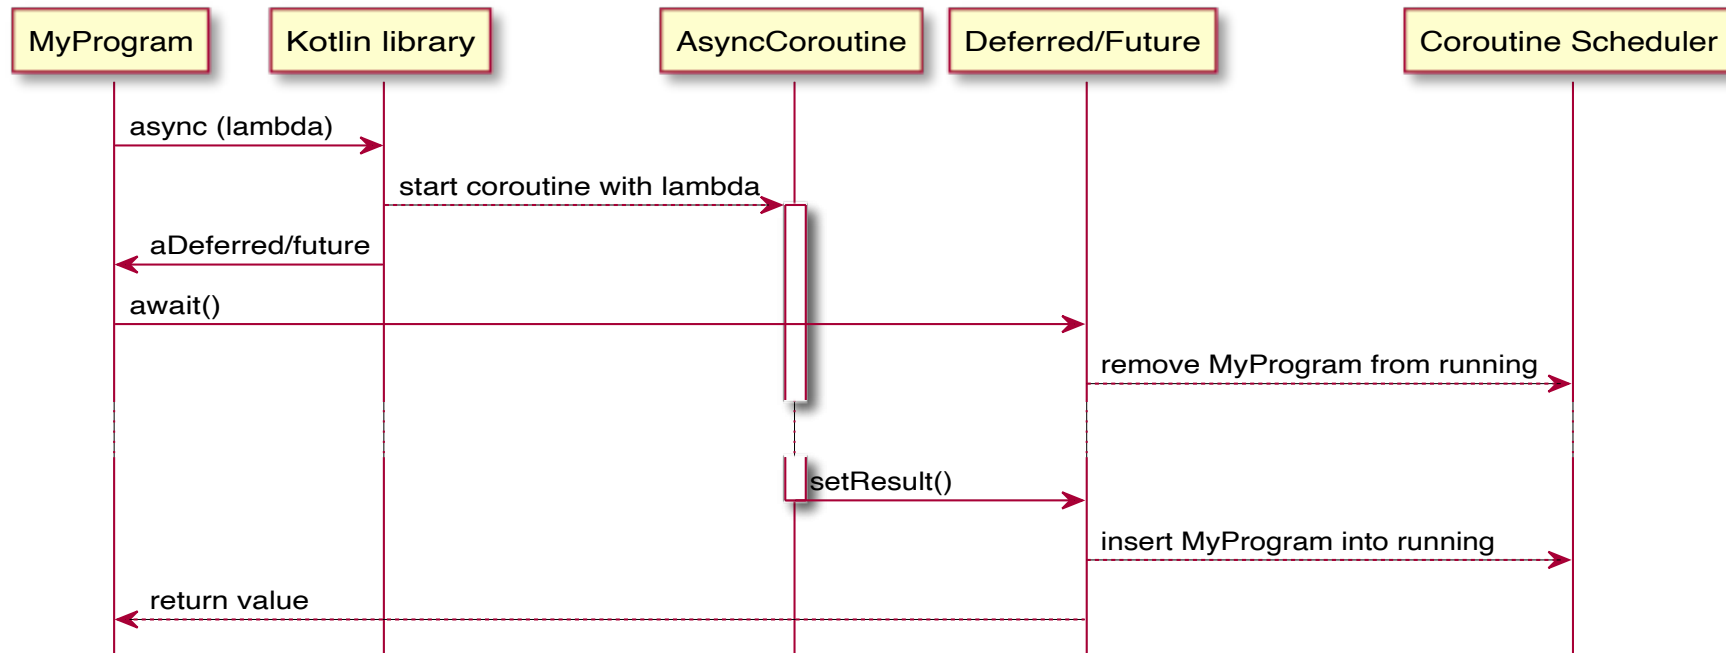
```
suspend fun produceNumbers(ch : Channel<Int>) {  
    var x = 1 // start from 1  
    while (true) {  
        ch.send(x++) // produce next  
        delay(100) // wait 0.1s  
    }  
}
```

```
suspend fun consumer(id: Int, channel: ReceiveChannel<Int>) {  
    for (msg in channel) { // reading from a channel  
        println("#$id --> $msg")  
    }  
}
```

Async & await

```
val one = async { computation1() }  
val two = async { computation2() }  
res = one.await() + two.await()
```

- The function `async` launches a coroutine,
- returns a `Deferred` (Kotlin name for future/promise)
- The `await()` method on a deferred waits for the result to be ready



Library defined scheduler

In some languages it is possible to write a scheduler in the language itself.

```
Scheduler sch = new Scheduler()  
sch.add( new Process() )  
sch.add( new Process() )  
sch.run()
```

Where Scheduler and Process are library defined classes (opposed to compiler defined).

Design options

1. Should `run` return a value
2. How should `Scheduler` work together with native `Threads`
3. Can you add `Process` after the scheduler has started?
4. Which *states* do a process have (running, runnable, waiting, terminated,...)
5. How do you specify a `Process` (using `Lambda` or subclassing for example)
6. Can `Scheduler` be preemptive, or do `Process` have to be cooperative

Mapping coroutines to JVM threads

```
fun main() = runBlocking<Unit> {  
    launch { printInfo("Parent    ") }  
    launch(Dispatchers.Unconfined) { printInfo("Unconfined") }  
    launch(Dispatchers.Default) { printInfo("Default    ") }  
    launch(newSingleThreadContext("MyOwnThread")) { printInfo("New        ") }  
}
```

Dispatchers are schedulers that run coroutines.

- launch() - context of the parent, main runBlocking coroutine
- launch(Dispatchers.**Unconfined**) - main thread
- launch(Dispatchers.**Default**) - a CachedThreadPool
- launch(newSingleThreadContext("MyOwnThread")) - will get its own new thread

[Link to code](#)

Move coroutines between schedulers

```
fun log(msg: String) = println("[${Thread.currentThread().name}] $msg")

fun main() {
    val ctx1 = newSingleThreadContext("Ctx1")
    val ctx2 = newSingleThreadContext("Ctx2")
    runBlocking(ctx1) {
        log("Started in ctx1")
        withContext(ctx2) {
            log("Working in ctx2")
        }
        log("Back to ctx1")
    }
}
```

[Code](#)

OK, it is possible, but can that have any possible use????

Thread confinement of access to shared data

No need to sync if only one thread modifies data

If coroutines are run in thread-pool with true parallelism, bad things™ might happen.

Non-synchronized updates - recap - saw this in first lesson

Sample from [documentation](#)

The fastest solution

(because this has an easy solution)

The scheduling solution

runBlocking

Notice that many places you see

```
fun main() = runBlocking { suspending lambda }
```

runBlocking is a function which runs a coroutine scheduler.

the suspending lambda is the primary coroutine it schedules

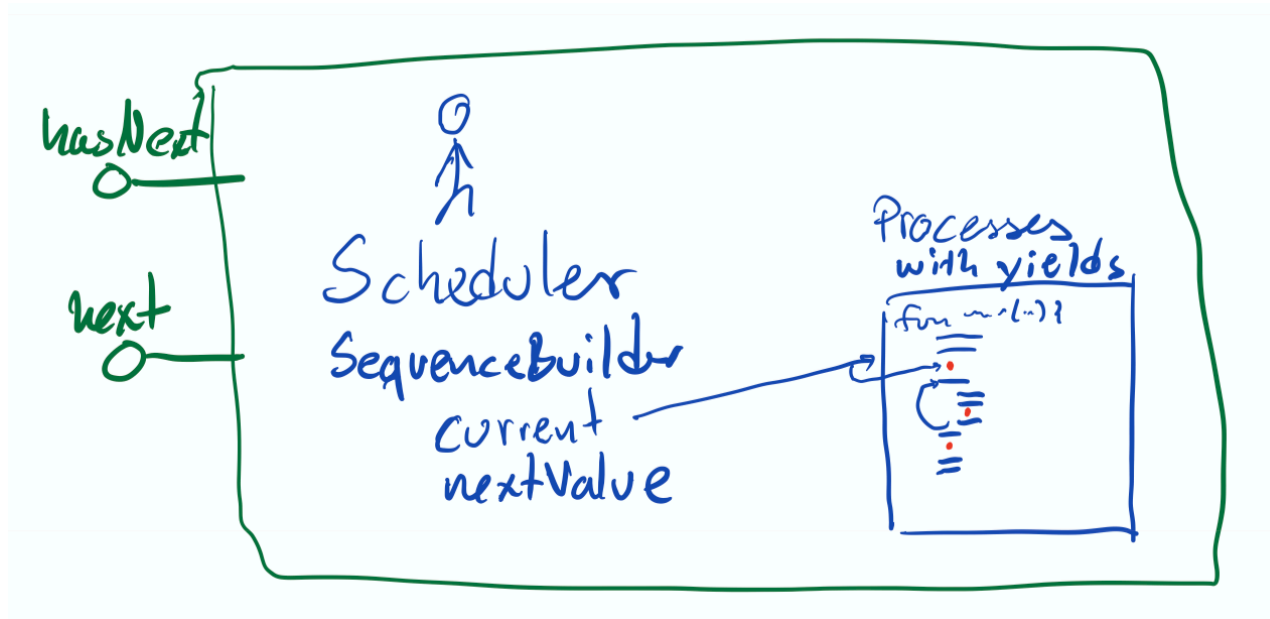
and if you spawn (launch or async) other coroutines from within, those run as well.

The scheduler runBlocking returns when **all** its scheduled coroutines have finished.

A (nearly) final word on yield

When the scheduler is an object (as opposed to a JVM concept), one can call methods on a scheduler.

The sequence function (where we use yield) is actually a scheduler with just one Process



Example of statemachine

```
fun foo() = sequence<Int> {  
    var i = 17  
    i -= 10  
    yield(i)  
    i += 18/9  
    yield(i)  
    i = 125/25+2*2*2  
    yield(i)  
}
```

called as : `foo().forEach(::println)`

There are two components here:

- the call to `sequence<In>`
- the *suspending lambda* which contains the yields

It is the suspending lambda that is translated into a state-machine

[Code in kotlin playground](#)

The suspending lambda is **translated** by the compiler into this

```
class FooMachine {
  var i = 0 // must be initialized to something
  var state = 0

  fun callMe():Int {
    var returnValue : Int
    when (state) {
      0 -> {i = 17
           i -= 10
           returnValue = i
           state = 1 }
      1 -> {i += 18/9
           returnValue = i
           state = 2 }
      2 -> {i = 125/25+2*2*2
           returnValue = i
           state = -1 }
      else -> throw IllegalStateException("Bad boy!")
    }
    return returnValue
  }
}
```

Finally the sequence function is

```
class FooIterator : Iterator<Int> {
    val machine = FooMachine()
    var nextValue : Int? = null

    init{ advance() }

    private fun advance(){
        if (machine.state != -1 )
            nextValue = machine.callMe()
        else
            nextValue = null
    }

    override fun hasNext():Boolean {
        return nextValue != null
    }

    override fun next(): Int {
        if ( nextValue == null) throw NoSuchElementException("Bad boy!")
        val ret : Int = nextValue!!
        advance()
        return ret
    }
}
```