

Practical Concurrent and Parallel Programming XIII

Lock free coordination

Jørgen Staunstrup, IT University

Friday 2020-11-20

Start at 8:00

1 / 30

Just for fun



Running TestCountFactors on my Android phone:
`final int range = 10000000;`

Results:



Samsung A20 - Lassen P+

4 cores

Time in s: 7.24

factors: 37861235



Lenovo ThinkBook Intel i5

8 cores (hyperthreading)

Time in s: 2.57

factors: 37861235

2 / 30

Apologies

Last week you asked questions about terminology:

concurrency, parallelism, multiprogramming, multithreading, multiprocessing, ...

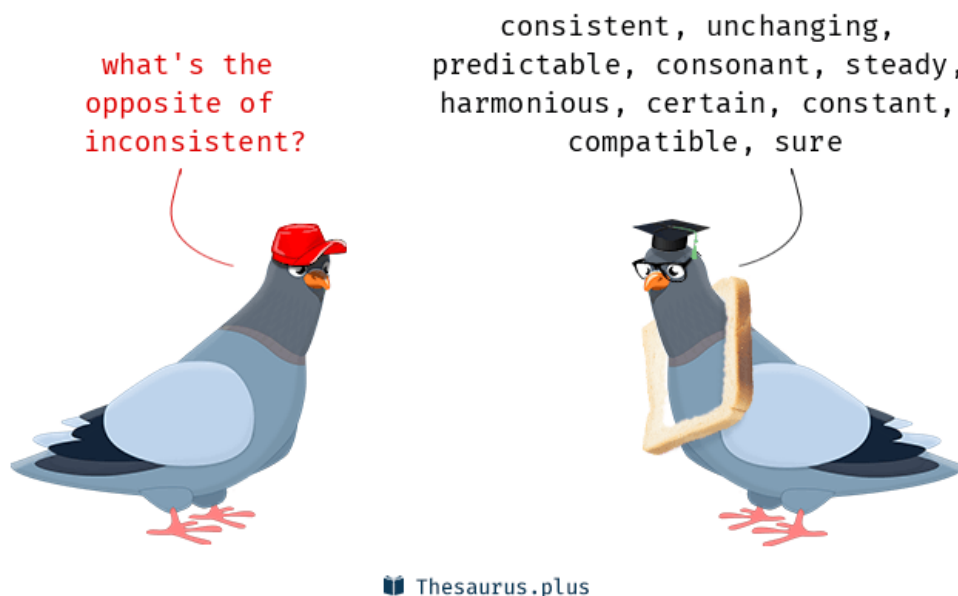
There is no commonly agreed consistent terminology !! The same is unfortunately the case for many other computer science topics.

What if the same was the case in medicine? elbow, armbender, armlink, ...

Learning and using the agreed terminology is a keep part of becoming a professional in many other fields

There is a similar terminology question to this week's material on LearnIT. I will return to this.

3 / 30



4 / 30

Agenda

Git

Optimistic concurrency control

Operational transform

Linearization

Realm (MongoDB) database

Atomicity

5 / 30

File sharing with Git

Branch: master ▾	PCPP / E2020 / Weeks /
jst final version of slides uploaded	
..	
01_BasicJava	issuing commit
02_BasicJava_II	week 2
03_Design	Start on 05, and some loose ends
04_PerformanceMeasurement	Start on 05, and some loose ends
05_Executor_Framework	Several changes
06_ParallelStreams	before fall break
07_Scaleability	some stuff from lesson 7
09_React	final version of slides uploaded
10_Coroutines-1	misc lesson 10 stuff

Workflow:

```
git pull % modifications
git stage -A
git commit ...
git push
```

Works because Kasper and I modify different files !!

6 / 30

Git merge / rebase

file abc.txt: abcdefg and file numbers.txt: 123456

GitEx: --all - gitk

File Edit View Help

● master 123456 and abcdefg

git branch newnumbers

git checkout newnumbers

change file numbers.txt: 1234

GitEx: --all - gitk

File Edit View Help

● newnumbers 1234
● master 123456 abcdefg

git checkout master

git merge newnumbers

Updating dd2289c..a423cf8

Fast-forward

numbers.txt | 2 +-
1 file changed, 1 insertion(+), 1 deletion(-)

7 / 30

Git merge/rebase (2)

GitEx: --all - gitk

File Edit View Help

● master 123456 and abcdefg

change file numbers.txt: 12xy4q

GitEx: --all - gitk

File Edit View Help

● newnumbers 12xy4q
● master 123456 abcdefg

git checkout master

git merge newnumbers

Auto-merging numbers.txt

CONFLICT (content): Merge conflict in numbers.txt

Automatic merge failed; fix conflicts and then commit the result.

8 / 30

Pessimistic concurrency control

```
public void synchronized modify(Something s) {  
    ...  
}
```

9 / 30

Optimistic concurrency control

```
public void ???? modify (Something s) {  
    ...  
}
```

Locking (atomicity) is pessimistic concurrency

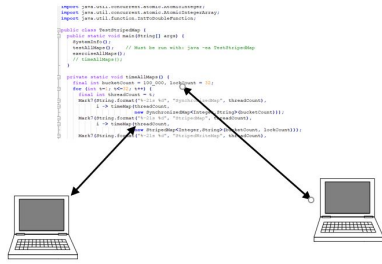
Google Wave and Realm (MongoDB) use optimistic concurrency control

Compromise on consistency: *Strong eventual consistency*

Recall discussion about `forEach` in exercise 7.2 (Strided map)

10 / 30

Concurrent text editing (Google wave)



Google wave <https://youtu.be/p6pgxLaDdQw>




Concurrent editing survived in Google Docs.

11 / 30

Operational transform

The key concept behind Google Wave (and many similar systems)



<https://youtu.be/3ykZYKCK7AM>


Key idea is to find a way to resolve conflicts for all pairs of operations $o1$ and $o2$ where: $o1;o2 \neq o2;o1$

This is not so difficult for text operations like insert and delete.


12 / 30

Consistency

Defining acceptable behavior of shared data.



Is this consistent?



Is this consistent?

13 / 30

LearnIT question

In the paper: Linearizability: A Correctness Condition for Concurrent Objects (WingHerlihyCons.pdf in this weeks reading) the authors write

Unlike alternative correctness conditions such as sequential consistency [31] or serializability [40], linearizability is a local property

Local property ???

14 / 30

Consistency definitions

Strict consistency: *a write to a variable by any thread needs to be seen instantaneously by all other threads.*

Example: a bank account

Note that for some datastructures strict consistency can be ensured even if some operations are not atomic.

Example?

Weak consistency: *A write to a variable does not have to be seen instantaneously, however, writes to variables by different threads have to be seen in the same order by all threads.*

15 / 30

Weak consistency: (example)

Example: the striped map from week 7

```
public void forEach(Consumer<K,V> consumer) {
    final ItemNode<K,V>[] bs = buckets;
    for (int stripe=0; stripe<lockCount; stripe++)
        synchronized (locks[stripe]) {
            for (int hash=stripe; hash<bs.length; hash+=lockCount) {
                ItemNode<K,V> node = bs[hash];
                while (node != null) {
                    consumer.accept(node.k, node.v);
                    node = node.next;
                }
            }
        }
}
```

t_1 : `forEach(bs[1])` — assume it takes a while

t_2 : `remove(bs[0])` and a little later `insert(bs[2])` — before t_1 moves on


Now t_1 will have iterated (`forEach`) over `bs[0]`, `bs[1]`, `bs[2]`, ...

!!!

Is this consistent?

16 / 30

Consistency definitions (2)



source: *Consistency in Non-Transactional Distributed Storage Systems* by Paolo Viotti and Marko Vukolic'

17 / 30

CAP theorem

Consistency

Every read receives the most recent write or an error

Availability

Every request receives a (non-error) response – without guarantee that it contains the most recent write

Partition tolerance

The system continues to operate despite an arbitrary number of messages being dropped (or delayed) by the network between nodes

CAP theorem: *impossible* for a distributed data store to simultaneously provide more than two out of the three: consistency, availability and partition tolerance.


Gilbert and Nancy Lynch, "Brewer's conjecture and the feasibility of consistent, available, partition-tolerant web services", ACM SIGACT News, Volume 33 Issue 2 (2002), pg. 51–59.
<https://dl.acm.org/doi/10.1145/564585.564601>

18 / 30

Strong eventual consistency

Off-line is default - AP system


When online, requests are merged (operational transform)



19 / 30

Operational transform (example)

Imagine a text editor where many clients can edit without locking



a shared document:

The server makes an opTrans operation on conflicting operations such as: del4 and del2.

```
opTrans(del x, del y) =  
  {delx-1, dely} if x>y  
  {delx, dely-1} if x<y  
  {no-op, no-op} if x = y
```

More details: *High-Latency, Low-Bandwidth Windowing in the Jupiter Collaboration System*, see Nichols.pdf In the folder with material for lecture 13).

But what if the operations are more complicated than insert and delete?

E.g. in a database?

20 / 30

Realm database

Mobile apps: network reliability, local storage, and UIs reactive.

- ▶ Local storage: on client devices i.e. fast.
- ▶ Network reliability: Realm Database is offline-first: you always read from and write to the local database. Synchronizes data with central database in a background thread. The sync protocol resolves conflicts using operational transform.
- ▶ Reactive UI: Live objects always reflect the latest data stored.
- ▶ Object oriented: Database stores Java objects directly.

The Realm SDK: Android, iOS, Node.js, React Native, and UWP (Windows).

Realm is now part of MongoDB.

source: <https://docs.mongodb.com/realm/get-started/introduction-mobile/>

21 / 30

Realm synchronization protocol

Goal: correctly and efficiently sync data changes in real time across multiple clients that each maintain their own local Realm files.

- ▶ Changeset: list of write operations to database objects.
- ▶ Operational transformation: operational transformation is used to resolve conflicts between changesets from different clients.
- ▶ Off-line first: any device may perform offline writes and upload the corresponding changesets when there is network connectivity.
- ▶ Realm objects: Some restrictions on field types (to enable operational transform).

source: <https://docs.mongodb.com/realm/sync/protocol/#sync-protocol>

22 / 30


Atomicity

Throughout the course we have relied on locks to ensure atomicity. Let us take a little closer look at how locks work.



23 / 30

How to implement atomicity



```
wait(s);  
    atomic operation;  
signal(s);
```

semaphore



24 / 30

How to implement a semaphore/lock on a computer

Assume that the computer hardware offers an atomic exchange operation: $a \leftrightarrow b$

```
boolean semaphore= true; //global
```

```
boolean enter= false; //local
```

```
repeat  
  enter  $\leftrightarrow$  semaphore  
until enter
```

```
atomic operation
```

```
enter  $\leftrightarrow$  semaphore
```

```
boolean enter= false; //local
```

```
repeat  
  enter  $\leftrightarrow$  semaphore  
until enter
```

...

```
atomic operation
```

```
enter  $\leftrightarrow$  semaphore
```

Modern hardware has such an operation e.g. on Intel `cmpxchg`

25 / 30

Multi Maren 1980




$a \leftrightarrow b$ implemented
by a hardware arbiter

26 / 30

Hardware implementation of $a \leftrightarrow b$

Gust professor at UW Washington 1984-85

Introduced to VLSI (chip) design



How to implement: $a \leftrightarrow b$
???

27 / 30

How to implement \leftrightarrow in hardware?

Clock: step1, step2, step3, ...

a single operation (e.g. $a \leftrightarrow b$) in each step !!!



BUT what about networked computers?

They have different clocks !!! same challenge with all other i/o



28 / 30

Arbiter



Anomalous Behavior of Synchronizer and Arbiter Circuits. Thomas J. Chaney and Charles E. Molnar, IEEE TC 22, April 1973

Buridan's donkey 1230

<https://en.wikipedia.org/wiki/Buridan>

https://en.wikipedia.org/wiki/Buridan%27s_ass

29 / 30

The end

This (philosophical) detour ends this course.

Next week: Info about examination and course evaluation

30 / 30